

How do I read a tabular data file into pandas?

```
In [1]: ▶ import pandas as pd
# read a dataset of Chipotle orders directly from a URL and store the results
orders = pd.read_table('http://bit.ly/chiporders')
```

```
In [2]: ▶ # examine the first 5 rows
orders.head()
```

Out[2]:

	order_id	quantity	item_name	choice_description	item_price
0	1	1	Chips and Fresh Tomato Salsa	NaN	\$2.39
1	1	1	Izze	[Clementine]	\$3.39
2	1	1	Nantucket Nectar	[Apple]	\$3.39
3	1	1	Chips and Tomatillo-Green Chili Salsa	NaN	\$2.39
4	2	2	Chicken Bowl	[Tomatillo-Red Chili Salsa (Hot), [Black Beans...	\$16.98

Documentation for `read_table` (http://pandas.pydata.org/pandas-docs/stable/generated/pandas.read_table.html).

```
In [3]: ▶ # read a dataset of movie reviewers (modifying the default parameter values of
user_cols = ['user_id', 'age', 'gender', 'occupation', 'zip_code']
user = pd.read_csv(r"user.txt", sep='|', header=None, names=user_cols)
```

```
In [4]: ▶ # csv - ',', tsv - \t, txt -
```

```
In [5]: ▶ # pd.read_excel('.xlsx', sheet_name='Sheet1')
```

```
In [6]: # examine the first 5 rows  
user.head(10)
```

Out[6]:

	user_id	age	gender	occupation	zip_code
0	1	24	M	technician	85711
1	2	53	F	other	94043
2	3	23	M	writer	32067
3	4	24	M	technician	43537
4	5	33	F	other	15213
5	6	42	M	executive	98101
6	7	57	M	administrator	91344
7	8	36	M	administrator	05201
8	9	29	M	student	01002
9	10	53	M	lawyer	90703

```
In [7]: # user.tail(10)
```

Out[7]:

	user_id	age	gender	occupation	zip_code
933	934	61	M	engineer	22902
934	935	42	M	doctor	66221
935	936	24	M	other	32789
936	937	48	M	educator	98072
937	938	38	F	technician	55038
938	939	26	F	student	33319
939	940	32	M	administrator	02215
940	941	20	M	student	97229
941	942	48	F	librarian	78209
942	943	22	M	student	77841

How do I select a pandas Series from a DataFrame?

```
In [8]: # read a dataset of UFO reports into a DataFrame  
ufo = pd.read_table('ufo.csv', sep=',')
```

```
In [9]: # read_csv is equivalent to read_table, except it assumes a comma separator  
ufo = pd.read_csv('ufo.csv')
```

```
In [10]: # examine the first 5 rows  
ufo.head()
```

Out[10]:

	City	Colors Reported	Shape Reported	State	Time
0	Ithaca	NaN	TRIANGLE	NY	6/1/1930 22:00
1	Willingboro	NaN	OTHER	NJ	6/30/1930 20:00
2	Holyoke	NaN	OVAL	CO	2/15/1931 14:00
3	Abilene	NaN	DISK	KS	6/1/1931 13:00
4	New York Worlds Fair	NaN	LIGHT	NY	4/18/1933 19:00

```
In [11]: # select the 'City' Series using bracket notation  
ufo['City']  
  
# or equivalently, use dot notation  
ufo.State
```

Out[11]:

0	NY
1	NJ
2	CO
3	KS
4	NY
	..
18236	IL
18237	IA
18238	WI
18239	WI
18240	FL

Name: State, Length: 18241, dtype: object

Bracket notation will always work, whereas **dot notation** has limitations:

- Dot notation doesn't work if there are **spaces** in the Series name
- Dot notation doesn't work if the Series has the same name as a **DataFrame method or attribute** (like 'head' or 'shape')
- Dot notation can't be used to define the name of a **new Series** (see below)

```
In [12]: ▶ # create a new 'Location' Series (must use bracket notation to define the Series)
ufo['Location'] = ufo.City + ', ' + ufo.State
ufo.head()
```

Out[12]:

	City	Colors Reported	Shape Reported	State	Time	Location
0	Ithaca	NaN	TRIANGLE	NY	6/1/1930 22:00	Ithaca, NY
1	Willingboro	NaN	OTHER	NJ	6/30/1930 20:00	Willingboro, NJ
2	Holyoke	NaN	OVAL	CO	2/15/1931 14:00	Holyoke, CO
3	Abilene	NaN	DISK	KS	6/1/1931 13:00	Abilene, KS
4	New York Worlds Fair	NaN	LIGHT	NY	4/18/1933 19:00	New York Worlds Fair, NY

```
In [13]: ▶ ufo.shape
```

Out[13]: (18241, 6)

Why do some pandas commands end with parentheses (and others don't)?

```
In [14]: ▶ # read a dataset of top-rated IMDb movies into a DataFrame
import pandas as pd
movies = pd.read_csv('http://bit.ly/imdbratings')
```

Methods end with parentheses, while **attributes** don't:

```
In [15]: ▶ # example method: show the first 5 rows
movies.head()
```

Out[15]:

	star_rating	title	content_rating	genre	duration	actors_list
0	9.3	The Shawshank Redemption	R	Crime	142	[u'Tim Robbins', u'Morgan Freeman', u'Bob Gunt...
1	9.2	The Godfather	R	Crime	175	[u'Marlon Brando', u'Al Pacino', u'James Caan']
2	9.1	The Godfather: Part II	R	Crime	200	[u'Al Pacino', u'Robert De Niro', u'Robert Duv...
3	9.0	The Dark Knight	PG-13	Action	152	[u'Christian Bale', u'Heath Ledger', u'Aaron E...
4	8.9	Pulp Fiction	R	Crime	154	[u'John Travolta', u'Uma Thurman', u'Samuel L....

```
In [16]: ▶ movies.genre.nunique()
```

Out[16]: 16

```
In [17]: ▶ movies.genre.unique()
```

Out[17]: array(['Crime', 'Action', 'Drama', 'Western', 'Adventure', 'Biography',
'Comedy', 'Animation', 'Mystery', 'Horror', 'Film-Noir', 'Sci-Fi',
'History', 'Thriller', 'Family', 'Fantasy'], dtype=object)

```
In [18]: ▶ # example method: calculate summary statistics
movies.describe()
```

Out[18]:

	star_rating	duration
count	979.000000	979.000000
mean	7.889785	120.979571
std	0.336069	26.218010
min	7.400000	64.000000
25%	7.600000	102.000000
50%	7.800000	117.000000
75%	8.100000	134.000000
max	9.300000	242.000000

In [19]: `movies.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 979 entries, 0 to 978
Data columns (total 6 columns):
star_rating      979 non-null float64
title            979 non-null object
content_rating   976 non-null object
genre            979 non-null object
duration         979 non-null int64
actors_list      979 non-null object
dtypes: float64(1), int64(1), object(4)
memory usage: 46.0+ KB
```

In [20]: `# example attribute: number of rows and columns`
`movies.shape`

Out[20]: (979, 6)

In [21]: `# example attribute: data type of each column`
`movies.dtypes`

Out[21]:

star_rating	float64
title	object
content_rating	object
genre	object
duration	int64
actors_list	object
dtype:	object

In [22]: `# use an optional parameter to the describe method to summarize only 'object'`
`movies.describe(include='all')`

Out[22]:

	star_rating	title	content_rating	genre	duration	actors_list
count	979.000000	979	976	979	979.000000	979
unique	NaN	975	12	16	NaN	969
top	NaN	Les Miserables	R	Drama	NaN	[u'Daniel Radcliffe', u'Emma Watson', u'Rupert...
freq	NaN	2	460	278	NaN	6
mean	7.889785	NaN	NaN	NaN	120.979571	NaN
std	0.336069	NaN	NaN	NaN	26.218010	NaN
min	7.400000	NaN	NaN	NaN	64.000000	NaN
25%	7.600000	NaN	NaN	NaN	102.000000	NaN
50%	7.800000	NaN	NaN	NaN	117.000000	NaN
75%	8.100000	NaN	NaN	NaN	134.000000	NaN
max	9.300000	NaN	NaN	NaN	242.000000	NaN

Documentation for **describe** [_\(http://pandas.pydata.org/pandas-docs/stable/generated/pandas.DataFrame.describe.html\)](http://pandas.pydata.org/pandas-docs/stable/generated/pandas.DataFrame.describe.html)

5. How do I rename columns in a pandas DataFrame?

```
In [23]:  ufo = pd.read_csv('ufo.csv')
```

```
In [24]:  # examine the column names
ufo.columns
```

```
Out[24]: Index(['City', 'Colors Reported', 'Shape Reported', 'State', 'Time'], dtype
='object')
```

```
In [25]:  # rename two of the columns by using the 'rename' method
ufo.rename(columns={'Colors Reported': 'Colors_Reported', 'Shape Reported': 'Shape_Reported'})
ufo.columns
```

```
Out[25]: Index(['City', 'Colors_Reported', 'Shape_Reported', 'State', 'Time'], dtype
='object')
```

Documentation for **rename** [_\(http://pandas.pydata.org/pandas-docs/stable/generated/pandas.DataFrame.rename.html\)](http://pandas.pydata.org/pandas-docs/stable/generated/pandas.DataFrame.rename.html)

```
In [26]:  ufo.columns
```

```
Out[26]: Index(['City', 'Colors_Reported', 'Shape_Reported', 'State', 'Time'], dtype
='object')
```

```
In [27]:  # replace all of the column names by overwriting the 'columns' attribute
ufo_cols = ['city', 'colors reported', 'shape reported', 'state', 'time']
ufo.columns = ufo_cols
ufo.columns
```

```
Out[27]: Index(['city', 'colors reported', 'shape reported', 'state', 'time'], dtype
='object')
```

```
In [28]:  # replace the column names during the file reading process by using the 'names' parameter
import pandas as pd
ufo = pd.read_csv('ufo.csv', header=0, names=ufo_cols)
ufo.columns
```

```
Out[28]: Index(['city', 'colors reported', 'shape reported', 'state', 'time'], dtype
='object')
```

Documentation for **read_csv** [_\(http://pandas.pydata.org/pandas-docs/stable/generated/pandas.read_csv.html\)](http://pandas.pydata.org/pandas-docs/stable/generated/pandas.read_csv.html)

```
In [29]: # replace all spaces with underscores in the column names by using the 'str.replace' method  
ufo.columns = ufo.columns.str.replace(' ', '_')  
ufo.columns
```

```
Out[29]: Index(['city', 'colors_reported', 'shape_reported', 'state', 'time'], dtype='object')
```

Documentation for `str.replace` (<http://pandas.pydata.org/pandas-docs/stable/generated/pandas.Series.str.replace.html>)

How do I remove columns from a pandas DataFrame?

```
In [30]: # read a dataset of UFO reports into a DataFrame  
ufo = pd.read_csv('ufo.csv')  
ufo.head()
```

```
Out[30]:
```

	City	Colors Reported	Shape Reported	State	Time
0	Ithaca	NaN	TRIANGLE	NY	6/1/1930 22:00
1	Willingboro	NaN	OTHER	NJ	6/30/1930 20:00
2	Holyoke	NaN	OVAL	CO	2/15/1931 14:00
3	Abilene	NaN	DISK	KS	6/1/1931 13:00
4	New York Worlds Fair	NaN	LIGHT	NY	4/18/1933 19:00

```
In [31]: # remove a single column (axis=1 refers to columns)  
ufo.drop('Colors Reported', axis=1, inplace=True)  
ufo.head()
```

```
Out[31]:
```

	City	Shape Reported	State	Time
0	Ithaca	TRIANGLE	NY	6/1/1930 22:00
1	Willingboro	OTHER	NJ	6/30/1930 20:00
2	Holyoke	OVAL	CO	2/15/1931 14:00
3	Abilene	DISK	KS	6/1/1931 13:00
4	New York Worlds Fair	LIGHT	NY	4/18/1933 19:00

Documentation for `drop` (<http://pandas.pydata.org/pandas-docs/stable/generated/pandas.DataFrame.drop.html>)


```
In [32]: # remove multiple columns at once
ufo.drop(['City', 'State'], axis=1)
ufo.head()
ufo.drop(['City', 'State'], axis=1).head()
```

Out[32]:

	Shape Reported	Time
0	TRIANGLE	6/1/1930 22:00
1	OTHER	6/30/1930 20:00
2	OVAL	2/15/1931 14:00
3	DISK	6/1/1931 13:00
4	LIGHT	4/18/1933 19:00

```
In [33]: ufo.head(5)
```

Out[33]:

	City	Shape Reported	State	Time
0	Ithaca	TRIANGLE	NY	6/1/1930 22:00
1	Willingboro	OTHER	NJ	6/30/1930 20:00
2	Holyoke	OVAL	CO	2/15/1931 14:00
3	Abilene	DISK	KS	6/1/1931 13:00
4	New York Worlds Fair	LIGHT	NY	4/18/1933 19:00

```
In [34]: ufo.reset_index(inplace=True)
```

```
In [35]: ufo.head()
```

Out[35]:

	index	City	Shape Reported	State	Time
0	0	Ithaca	TRIANGLE	NY	6/1/1930 22:00
1	1	Willingboro	OTHER	NJ	6/30/1930 20:00
2	2	Holyoke	OVAL	CO	2/15/1931 14:00
3	3	Abilene	DISK	KS	6/1/1931 13:00
4	4	New York Worlds Fair	LIGHT	NY	4/18/1933 19:00

```
In [36]: # remove multiple rows at once (axis=0 refers to rows)
ufo.drop([2,9], axis=0, inplace=True)
# ufo.reset_index()
```

How do I sort a pandas DataFrame or a Series?

```
In [37]: # read a dataset of top-rated IMDb movies into a DataFrame
import pandas as pd
movies = pd.read_csv('http://bit.ly/imdbratings')
movies.head()
```

Out[37]:

	star_rating	title	content_rating	genre	duration	actors_list
0	9.3	The Shawshank Redemption	R	Crime	142	[u'Tim Robbins', u'Morgan Freeman', u'Bob Gunt...
1	9.2	The Godfather	R	Crime	175	[u'Marlon Brando', u'Al Pacino', u'James Caan']
2	9.1	The Godfather: Part II	R	Crime	200	[u'Al Pacino', u'Robert De Niro', u'Robert Duv...
3	9.0	The Dark Knight	PG-13	Action	152	[u'Christian Bale', u'Heath Ledger', u'Aaron E...
4	8.9	Pulp Fiction	R	Crime	154	[u'John Travolta', u'Uma Thurman', u'Samuel L....

Note: None of the sorting methods below affect the underlying data. (In other words, the sorting is temporary).

```
In [38]: movies.actors_list.apply(pd.Series)
```

Out[38]:

	0
0	[u'Tim Robbins', u'Morgan Freeman', u'Bob Gunt...
1	[u'Marlon Brando', u'Al Pacino', u'James Caan']
2	[u'Al Pacino', u'Robert De Niro', u'Robert Duv...
3	[u'Christian Bale', u'Heath Ledger', u'Aaron E...
4	[u'John Travolta', u'Uma Thurman', u'Samuel L....
...	...
974	[u'Dustin Hoffman', u'Jessica Lange', u'Teri G...
975	[u'Michael J. Fox', u'Christopher Lloyd', u'Ma...
976	[u'Russell Crowe', u'Paul Bettany', u'Billy Bo...
977	[u'JoBeth Williams', u'Heather O'Rourke', u'Cr...
978	[u'Charlie Sheen', u'Michael Douglas', u'Tamar...

979 rows × 1 columns

```
In [39]: # sort the 'title' Series in ascending order (returns a Series)
movies.star_rating.sort_values().head()
```

```
Out[39]: 978    7.4
          950    7.4
          949    7.4
          948    7.4
          947    7.4
          Name: star_rating, dtype: float64
```

```
In [40]: # sort in descending order instead
movies.star_rating.sort_values(ascending=False).head()
```


```
Out[40]: 0    9.3
          1    9.2
          2    9.1
          3    9.0
          6    8.9
          Name: star_rating, dtype: float64
```

Documentation for **sort_values** (http://pandas.pydata.org/pandas-docs/stable/generated/pandas.Series.sort_values.html) for a **Series**. (Prior to version 0.17, use **order** (<http://pandas.pydata.org/pandas-docs/version/0.17.0/generated/pandas.Series.order.html>), instead.)

```
In [41]: # sort the entire DataFrame by the 'title' Series (returns a DataFrame)
movies.sort_values('title').head()
```

```
Out[41]:
```

	star_rating	title	content_rating	genre	duration	actors_list
542	7.8	(500) Days of Summer	PG-13	Comedy	95	[u'Zooey Deschanel', u'Joseph Gordon-Levitt', ...]
5	8.9	12 Angry Men	NOT RATED	Drama	96	[u'Henry Fonda', u'Lee J. Cobb', u'Martin Bals...]
201	8.1	12 Years a Slave	R	Biography	134	[u'Chiwetel Ejiofor', u'Michael Kenneth Willia...]
698	7.6	127 Hours	R	Adventure	94	[u'James Franco', u'Amber Tamblyn', u'Kate Mara']
110	8.3	2001: A Space Odyssey	G	Mystery	160	[u'Keir Dullea', u'Gary Lockwood', u'William S...]

```
In [42]:  # sort in descending order instead
movies.sort_values('title', ascending=False).head()
```

Out[42]:

	star_rating	title	content_rating	genre	duration	actors_list
864	7.5	[Rec]	R	Horror	78	[u'Manuela Velasco', u'Ferran Terraza', u'Jorg...
526	7.8	Zulu	UNRATED	Drama	138	[u'Stanley Baker', u'Jack Hawkins', u'Ulla Jac...
615	7.7	Zombieland	R	Comedy	88	[u'Jesse Eisenberg', u'Emma Stone', u'Woody Ha...
677	7.7	Zodiac	R	Crime	157	[u'Jake Gyllenhaal', u'Robert Downey Jr.', u'M...
955	7.4	Zero Dark Thirty	R	Drama	157	[u'Jessica Chastain', u'Joel Edgerton', u'Chri...

Documentation for **sort_values** (http://pandas.pydata.org/pandas-docs/stable/generated/pandas.DataFrame.sort_values.html) for a **DataFrame**. (Prior to version 0.17, use **sort** (<http://pandas.pydata.org/pandas-docs/version/0.17.0/generated/pandas.DataFrame.sort.html>) instead.)

```
In [43]: # sort the DataFrame first by 'content_rating', then by 'duration'
movies.sort_values(['star_rating', 'duration', 'content_rating'], ascending=[False, True, True])
```

Out[43]:

	star_rating	title	content_rating	genre	duration	actors_list
0	9.3	The Shawshank Redemption	R	Crime	142	[u'Tim Robbins', u'Morgan Freeman', u'Bob Gunt...
1	9.2	The Godfather	R	Crime	175	[u'Marlon Brando', u'Al Pacino', u'James Caan']
2	9.1	The Godfather: Part II	R	Crime	200	[u'Al Pacino', u'Robert De Niro', u'Robert Duv...
3	9.0	The Dark Knight	PG-13	Action	152	[u'Christian Bale', u'Heath Ledger', u'Aaron E...
5	8.9	12 Angry Men	NOT RATED	Drama	96	[u'Henry Fonda', u'Lee J. Cobb', u'Martin Bals...
9	8.9	Fight Club	R	Drama	139	[u'Brad Pitt', u'Edward Norton', u'Helena Bonh...
4	8.9	Pulp Fiction	R	Crime	154	[u'John Travolta', u'Uma Thurman', u'Samuel L....
6	8.9	The Good, the Bad and the Ugly	NOT RATED	Western	161	[u'Clint Eastwood', u'Eli Wallach', u'Lee Van ...
8	8.9	Schindler's List	R	Biography	195	[u'Liam Neeson', u'Ralph Fiennes', u'Ben Kings...
7	8.9	The Lord of the Rings: The Return of the King	PG-13	Adventure	201	[u'Elijah Wood', u'Viggo Mortensen', u'Ian McK...

How do I filter rows of a pandas DataFrame by column value?

```
In [44]: # read a dataset of top-rated IMDb movies into a DataFrame
movies = pd.read_csv('http://bit.ly/imdbratings')
movies.head()
```

Out[44]:

	star_rating	title	content_rating	genre	duration	actors_list
0	9.3	The Shawshank Redemption	R	Crime	142	[u'Tim Robbins', u'Morgan Freeman', u'Bob Gunt...
1	9.2	The Godfather	R	Crime	175	[u'Marlon Brando', u'Al Pacino', u'James Caan']
2	9.1	The Godfather: Part II	R	Crime	200	[u'Al Pacino', u'Robert De Niro', u'Robert Duv...
3	9.0	The Dark Knight	PG-13	Action	152	[u'Christian Bale', u'Heath Ledger', u'Aaron E...
4	8.9	Pulp Fiction	R	Crime	154	[u'John Travolta', u'Uma Thurman', u'Samuel L....

```
In [45]: # examine the number of rows and columns
movies.shape
```

Out[45]: (979, 6)

Goal: Filter the DataFrame rows to only show movies with a 'duration' of at least 200 minutes.

```
In [46]: movies.duration
```

Out[46]:

```
0      142
1      175
2      200
3      152
4      154
...
974    116
975    118
976    138
977    114
978    126
Name: duration, Length: 979, dtype: int64
```

```
In [47]: # create a list in which each element refers to a DataFrame row: True if the
booleans = []
for length in movies.duration:
    if length >= 200:
        booleans.append(True)
    else:
        booleans.append(False)
```

```
In [48]: ▶ # confirm that the list has the same length as the DataFrame  
len(booleans)
```

Out[48]: 979

```
In [49]: ▶ # examine the first five list elements  
booleans[0:5]
```

Out[49]: [False, False, True, False, False]

```
In [50]: ▶ # convert the list to a Series  
is_long = pd.Series(booleans)  
is_long.head(10)
```

Out[50]:

0	False
1	False
2	True
3	False
4	False
5	False
6	False
7	True
8	False
9	False

dtype: bool

```
In [51]: # use bracket notation with the boolean Series to tell the DataFrame which rows are long
movies[movies['is_long']]
```

Out[51]:

star_rating		title	content_rating	genre	duration	actors_list
2	9.1	The Godfather: Part II	R	Crime	200	[u'Al Pacino', u'Robert De Niro', u'Robert Duv...
7	8.9	The Lord of the Rings: The Return of the King	PG-13	Adventure	201	[u'Elijah Wood', u'Viggo Mortensen', u'Ian McK...
17	8.7	Seven Samurai	UNRATED	Drama	207	[u'Toshirō Mifune', u'Takashi Shimura', u'K...
78	8.4	Once Upon a Time in America	R	Crime	229	[u'Robert De Niro', u'James Woods', u'Elizabet...
85	8.4	Lawrence of Arabia	PG	Adventure	216	[u"Peter O'Toole", u'Alec Guinness', u'Anthony...
142	8.3	Lagaan: Once Upon a Time in India	PG	Adventure	224	[u'Aamir Khan', u'Gracy Singh', u'Rachel Shell...
157	8.2	Gone with the Wind	G	Drama	238	[u'Clark Gable', u'Vivien Leigh', u'Thomas Mit...
204	8.1	Ben-Hur	G	Adventure	212	[u'Charlton Heston', u'Jack Hawkins', u'Stephe...
445	7.9	The Ten Commandments	APPROVED	Adventure	220	[u'Charlton Heston', u'Yul Brynner', u'Anne Ba...
476	7.8	Hamlet	PG-13	Drama	242	[u'Kenneth Branagh', u'Julie Christie', u'Dere...
630	7.7	Malcolm X	PG-13	Biography	202	[u'Denzel Washington', u'Angela Bassett', u'De...

	star_rating	title	content_rating	genre	duration	actors_list
767	7.6	It's a Mad, Mad, Mad, Mad World	APPROVED	Action	205	[u'Spencer Tracy', u'Milton Berle', u'Ethel Me...

In [52]: `movies.duration >= 200`

```
Out[52]: 0      False
          1      False
          2       True
          3      False
          4      False
          ...
          974    False
          975    False
          976    False
          977    False
          978    False
          Name: duration, Length: 979, dtype: bool
```

```
In [53]: movies[movies.duration >= 200]
```

Out[53]:

star_rating		title	content_rating	genre	duration	actors_list
2	9.1	The Godfather: Part II	R	Crime	200	[u'Al Pacino', u'Robert De Niro', u'Robert Duv...
7	8.9	The Lord of the Rings: The Return of the King	PG-13	Adventure	201	[u'Elijah Wood', u'Viggo Mortensen', u'Ian McK...
17	8.7	Seven Samurai	UNRATED	Drama	207	[u'Toshirō Mifune', u'Takashi Shimura', u'K...
78	8.4	Once Upon a Time in America	R	Crime	229	[u'Robert De Niro', u'James Woods', u'Elizabet...
85	8.4	Lawrence of Arabia	PG	Adventure	216	[u"Peter O'Toole", u'Alec Guinness', u'Anthony...
142	8.3	Lagaan: Once Upon a Time in India	PG	Adventure	224	[u'Aamir Khan', u'Gracy Singh', u'Rachel Shell...
157	8.2	Gone with the Wind	G	Drama	238	[u'Clark Gable', u'Vivien Leigh', u'Thomas Mit...
204	8.1	Ben-Hur	G	Adventure	212	[u'Charlton Heston', u'Jack Hawkins', u'Stephe...
445	7.9	The Ten Commandments	APPROVED	Adventure	220	[u'Charlton Heston', u'Yul Brynner', u'Anne Ba...
476	7.8	Hamlet	PG-13	Drama	242	[u'Kenneth Branagh', u'Julie Christie', u'Dere...
630	7.7	Malcolm X	PG-13	Biography	202	[u'Denzel Washington', u'Angela Bassett', u'De...

	star_rating	title	content_rating	genre	duration	actors_list
767	7.6	It's a Mad, Mad, Mad, Mad World	APPROVED	Action	205	[u'Spencer Tracy', u'Milton Berle', u'Ethel Me...

```
In [54]: # simplify the steps above: no need to write a for loop to create 'is_long' s
is_long = movies.duration >= 200
movies[is_long]

# or equivalently, write it in one line (no need to create the 'is_long' obje
movies[movies.duration >= 200]
```

Out[54]:

	star_rating	title	content_rating	genre	duration	actors_list
2	9.1	The Godfather: Part II	R	Crime	200	[u'Al Pacino', u'Robert De Niro', u'Robert Duv...
7	8.9	The Lord of the Rings: The Return of the King	PG-13	Adventure	201	[u'Elijah Wood', u'Viggo Mortensen', u'Ian McK...
17	8.7	Seven Samurai	UNRATED	Drama	207	[u'Toshirxf4 Mifune', u'Takashi Shimura', u'K...
78	8.4	Once Upon a Time in America	R	Crime	229	[u'Robert De Niro', u'James Woods', u'Elizabet...
85	8.4	Lawrence of Arabia	PG	Adventure	216	[u"Peter O'Toole", u'Alec Guinness', u'Anthony...
142	8.3	Lagaan: Once Upon a Time in India	PG	Adventure	224	[u'Aamir Khan', u'Gracy Singh', u'Rachel Shell...
157	8.2	Gone with the Wind	G	Drama	238	[u'Clark Gable', u'Vivien Leigh', u'Thomas Mit...
204	8.1	Ben-Hur	G	Adventure	212	[u'Charlton Heston', u'Jack Hawkins', u'Stephe...
445	7.9	The Ten Commandments	APPROVED	Adventure	220	[u'Charlton Heston', u'Yul Brynnner', u'Anne Ba...
476	7.8	Hamlet	PG-13	Drama	242	[u'Kenneth Branagh', u'Julie Christie', u'Dere...

	star_rating	title	content_rating	genre	duration	actors_list
630	7.7	Malcolm X	PG-13	Biography	202	[u'Denzel Washington', u'Angela Bassett', u'De...]
767	7.6	It's a Mad, Mad, Mad, Mad World	APPROVED	Action	205	[u'Spencer Tracy', u'Milton Berle', u'Ethel Me...]

```
In [55]: # select the 'genre' Series from the filtered DataFrame
movies[movies.duration >= 200].genre

# or equivalently, use the 'loc' method
movies.loc[0:2, 'star_rating'] # t OR f OR INDEX , COLUMN NAME
```

```
Out[55]: 0    9.3
         1    9.2
         2    9.1
         Name: star_rating, dtype: float64
```

Documentation for `loc` (<http://pandas.pydata.org/pandas-docs/stable/generated/pandas.DataFrame.loc.html>)

How do I apply multiple filter criteria to a pandas DataFrame?

```
In [56]: # read a dataset of top-rated IMDb movies into a DataFrame
movies = pd.read_csv('http://bit.ly/imdbratings')
movies.head()
```

```
Out[56]:
```

	star_rating	title	content_rating	genre	duration	actors_list
0	9.3	The Shawshank Redemption	R	Crime	142	[u'Tim Robbins', u'Morgan Freeman', u'Bob Gunt...]
1	9.2	The Godfather	R	Crime	175	[u'Marlon Brando', u'Al Pacino', u'James Caan']
2	9.1	The Godfather: Part II	R	Crime	200	[u'Al Pacino', u'Robert De Niro', u'Robert Duv...]
3	9.0	The Dark Knight	PG-13	Action	152	[u'Christian Bale', u'Heath Ledger', u'Aaron E...]
4	8.9	Pulp Fiction	R	Crime	154	[u'John Travolta', u'Uma Thurman', u'Samuel L....]

In [57]: `# filter the DataFrame to only show movies with a 'duration' of at least 200`
`movies[movies.duration >= 200]`


Out[57]:

	star_rating	title	content_rating	genre	duration	actors_list
2	9.1	The Godfather: Part II	R	Crime	200	[u'Al Pacino', u'Robert De Niro', u'Robert Duv...
7	8.9	The Lord of the Rings: The Return of the King	PG-13	Adventure	201	[u'Elijah Wood', u'Viggo Mortensen', u'Ian McK...
17	8.7	Seven Samurai	UNRATED	Drama	207	[u'Toshirō Mifune', u'Takashi Shimura', u'K...
78	8.4	Once Upon a Time in America	R	Crime	229	[u'Robert De Niro', u'James Woods', u'Elizabet...
85	8.4	Lawrence of Arabia	PG	Adventure	216	[u"Peter O'Toole", u'Alec Guinness', u'Anthony...
142	8.3	Lagaan: Once Upon a Time in India	PG	Adventure	224	[u'Aamir Khan', u'Gracy Singh', u'Rachel Shell...
157	8.2	Gone with the Wind	G	Drama	238	[u'Clark Gable', u'Vivien Leigh', u'Thomas Mit...
204	8.1	Ben-Hur	G	Adventure	212	[u'Charlton Heston', u'Jack Hawkins', u'Stephe...
445	7.9	The Ten Commandments	APPROVED	Adventure	220	[u'Charlton Heston', u'Yul Brynner', u'Anne Ba...
476	7.8	Hamlet	PG-13	Drama	242	[u'Kenneth Branagh', u'Julie Christie', u'Dere...
630	7.7	Malcolm X	PG-13	Biography	202	[u'Denzel Washington', u'Angela Bassett', u'De...

	star_rating	title	content_rating	genre	duration	actors_list
767	7.6	It's a Mad, Mad, Mad, Mad World	APPROVED	Action	205	[u'Spencer Tracy', u'Milton Berle', u'Ethel Me...


Understanding **logical operators**:

- **and** : True only if **both sides** of the operator are True
- **or** : True if **either side** of the operator is True

In [58]:  *# demonstration of the 'and' operator*

```
print(True and True)
print(True and False)
print(False and False)
```

```
True
False
False
```

In [59]:  *# demonstration of the 'or' operator*

```
print(True or True)
print(True or False)
print(False or False)
```

```
True
True
False
```

Rules for specifying **multiple filter criteria** in pandas:

- use **&** instead of **and**
- use **|** instead of **or**
- add **parentheses** around each condition to specify evaluation order

Goal: Further filter the DataFrame of long movies (duration >= 200) to only show movies which also have a 'genre' of 'Drama'

```
In [60]: # CORRECT: use the '&' operator to specify that both conditions are required
movies[(movies.duration >=200) & (movies.genre == 'Drama')]
```

Out[60]:

	star_rating	title	content_rating	genre	duration	actors_list
17	8.7	Seven Samurai	UNRATED	Drama	207	[u'Toshir�xf4 Mifune', u'Takashi Shimura', u'K...
157	8.2	Gone with the Wind	G	Drama	238	[u'Clark Gable', u'Vivien Leigh', u'Thomas Mit...
476	7.8	Hamlet	PG-13	Drama	242	[u'Kenneth Branagh', u'Julie Christie', u'Dere...

```
In [61]: # INCORRECT: using the '|' operator would have shown movies that are either
movies[(movies.duration >=200) | (movies.genre == 'Drama')].head()
```

Out[61]:

	star_rating	title	content_rating	genre	duration	actors_list
2	9.1	The Godfather: Part II	R	Crime	200	[u'Al Pacino', u'Robert De Niro', u'Robert Duv...
5	8.9	12 Angry Men	NOT RATED	Drama	96	[u'Henry Fonda', u'Lee J. Cobb', u'Martin Bals...
7	8.9	The Lord of the Rings: The Return of the King	PG-13	Adventure	201	[u'Elijah Wood', u'Viggo Mortensen', u'Ian McK...
9	8.9	Fight Club	R	Drama	139	[u'Brad Pitt', u'Edward Norton', u'Helena Bonh...
13	8.8	Forrest Gump	PG-13	Drama	142	[u'Tom Hanks', u'Robin Wright', u'Gary Sinise']

Goal: Filter the original DataFrame to show movies with a 'genre' of 'Crime' or 'Drama' or 'Action'


```
In [62]: # use the '/' operator to specify that a row can match any of the three criteria
movies[(movies.genre == 'Crime') | (movies.genre == 'Drama') | (movies.genre == 'Action')]

# or equivalently, use the 'isin' method
movies[movies.genre.isin(['Crime', 'Drama', 'Action'])].head(10)

movies[movies.genre.isin(['Crime', 'Action', 'Drama'])]
```

Out[62]:

	star_rating	title	content_rating	genre	duration	actors_list
0	9.3	The Shawshank Redemption	R	Crime	142	[u'Tim Robbins', u'Morgan Freeman', u'Bob Gunt...]
1	9.2	The Godfather	R	Crime	175	[u'Marlon Brando', u'Al Pacino', u'James Caan']
2	9.1	The Godfather: Part II	R	Crime	200	[u'Al Pacino', u'Robert De Niro', u'Robert Duv...]
3	9.0	The Dark Knight	PG-13	Action	152	[u'Christian Bale', u'Heath Ledger', u'Aaron E...]
4	8.9	Pulp Fiction	R	Crime	154	[u'John Travolta', u'Uma Thurman', u'Samuel L....]
...
970	7.4	Wonder Boys	R	Drama	107	[u'Michael Douglas', u'Tobey Maguire', u'Franc...]
972	7.4	Blue Valentine	NC-17	Drama	112	[u'Ryan Gosling', u'Michelle Williams', u'John...]
973	7.4	The Cider House Rules	PG-13	Drama	126	[u'Tobey Maguire', u'Charlize Theron', u'Micha...]
976	7.4	Master and Commander: The Far Side of the World	PG-13	Action	138	[u'Russell Crowe', u'Paul Bettany', u'Billy Bo...]
978	7.4	Wall Street	R	Crime	126	[u'Charlie Sheen', u'Michael Douglas', u'Tamar...]

538 rows × 6 columns

Documentation for `isin` (<http://pandas.pydata.org/pandas-docs/stable/generated/pandas.Series.isin.html>)

When reading from a file, how do I read in only a subset of the columns?

```
In [63]: ▶ # read a dataset of UFO reports into a DataFrame, and check the columns
import pandas as pd
ufo = pd.read_csv('ufo.csv')
ufo.columns
```

```
Out[63]: Index(['City', 'Colors Reported', 'Shape Reported', 'State', 'Time'], dtype='object')
```

```
In [64]: ▶ # specify which columns to include by name
ufo = pd.read_csv('ufo.csv', usecols=['City', 'State'])
print(ufo.columns)
# or equivalently, specify columns by position
ufo = pd.read_csv('ufo.csv', usecols=[0, 4])
ufo.columns
```

```
Index(['City', 'State'], dtype='object')
```

```
Out[64]: Index(['City', 'Time'], dtype='object')
```

Question: When reading from a file, how do I read in only a subset of the rows?

```
In [65]: ▶ # specify how many rows to read
ufo = pd.read_csv('ufo.csv', nrows=3)
ufo
```

```
Out[65]:
```

	City	Colors Reported	Shape Reported	State	Time
0	Ithaca	NaN	TRIANGLE	NY	6/1/1930 22:00
1	Willingboro	NaN	OTHER	NJ	6/30/1930 20:00
2	Holyoke	NaN	OVAL	CO	2/15/1931 14:00

Documentation for `read_csv` (http://pandas.pydata.org/pandas-docs/stable/generated/pandas.read_csv.html)

Documentation for `select_dtypes` (http://pandas.pydata.org/pandas-docs/stable/generated/pandas.DataFrame.select_dtypes.html)

How do I use string methods in pandas?

```
In [66]: # read a dataset of Chipotle orders into a DataFrame
orders = pd.read_table('http://bit.ly/chiporders')
orders.head()
```

Out[66]:

	order_id	quantity	item_name	choice_description	item_price
0	1	1	Chips and Fresh Tomato Salsa	NaN	\$2.39
1	1	1	Izze	[Clementine]	\$3.39
2	1	1	Nantucket Nectar	[Apple]	\$3.39
3	1	1	Chips and Tomatillo-Green Chili Salsa	NaN	\$2.39
4	2	2	Chicken Bowl	[Tomatillo-Red Chili Salsa (Hot), [Black Beans...	\$16.98

```
In [67]: # normal way to access string methods in Python
'hello'.upper()
```

Out[67]: 'HELLO'

```
In [68]: # string methods for pandas Series are accessed via 'str'
orders.item_name.str.upper().head()
```

Out[68]:

```
0      CHIPS AND FRESH TOMATO SALSA
1                        IZZE
2      NANTUCKET NECTAR
3  CHIPS AND TOMATILLO-GREEN CHILI SALSA
4      CHICKEN BOWL
Name: item_name, dtype: object
```

```
In [69]: # string method 'contains' checks for a substring and returns a boolean Series
orders.item_name.str.contains('Chicken').head()
```

Out[69]:

```
0    False
1    False
2    False
3    False
4     True
Name: item_name, dtype: bool
```

```
In [70]: # use the boolean Series to filter the DataFrame
orders[orders.item_name.str.contains('Chicken')].head()
```

Out[70]:

	order_id	quantity	item_name	choice_description	item_price
	4	2	2	Chicken Bowl [Tomatillo-Red Chili Salsa (Hot), [Black Beans...	\$16.98
	5	3	1	Chicken Bowl [Fresh Tomato Salsa (Mild), [Rice, Cheese, Sou...	\$10.98
	11	6	1	Chicken Crispy Tacos [Roasted Chili Corn Salsa, [Fajita Vegetables,...	\$8.75
	12	6	1	Chicken Soft Tacos [Roasted Chili Corn Salsa, [Rice, Black Beans,...	\$8.75
	13	7	1	Chicken Bowl [Fresh Tomato Salsa, [Fajita Vegetables, Rice,...	\$11.25

```
In [71]: # string methods can be chained together
orders.choice_description.str.replace('[', '').str.replace(']', '').str.repla
```

Out[71]:

```
0      NaN
1    Clementine
2      Apple
3      NaN
4  Tomatillo-Red Chili Salsa (Hot) Black Beans Ri...
Name: choice_description, dtype: object
```

```
In [72]: # many pandas string methods support regular expressions (regex)
orders.choice_description.str.replace('[\[\]]', '').head()
```

Out[72]:

```
0      NaN
1    Clementine
2      Apple
3      NaN
4  Tomatillo-Red Chili Salsa (Hot), Black Beans, ...
Name: choice_description, dtype: object
```

[String handling section \(http://pandas.pydata.org/pandas-docs/stable/api.html#string-handling\)](http://pandas.pydata.org/pandas-docs/stable/api.html#string-handling) of the pandas API reference

How do I change the data type of a pandas Series?

```
In [73]: # read a dataset of alcohol consumption into a DataFrame  
drinks = pd.read_csv('http://bit.ly/drinksbycountry')  
drinks.head()
```

Out[73]:

	country	beer_servings	spirit_servings	wine_servings	total_litres_of_pure_alcohol	continent
0	Afghanistan	0	0	0	0.0	
1	Albania	89	132	54	4.9	Eu
2	Algeria	25	0	14	0.7	A
3	Andorra	245	138	312	12.4	Eu
4	Angola	217	57	45	5.9	A

```
In [74]: # examine the data type of each Series  
drinks.dtypes
```

Out[74]:

country	object
beer_servings	int64
spirit_servings	int64
wine_servings	int64
total_litres_of_pure_alcohol	float64
continent	object
dtype:	object

```
In [75]: # change the data type of an existing Series  
drinks['beer_servings'] = drinks.beer_servings.astype(float)  
drinks.dtypes
```

Out[75]:

country	object
beer_servings	float64
spirit_servings	int64
wine_servings	int64
total_litres_of_pure_alcohol	float64
continent	object
dtype:	object

Documentation for **astype** (<http://pandas.pydata.org/pandas-docs/stable/generated/pandas.Series.astype.html>)

```
In [76]: # alternatively, change the data type of a Series while reading in a file
drinks = pd.read_csv('http://bit.ly/drinksbycountry', dtype={'beer_servings':
drinks.dtypes
```

```
Out[76]: country                object
beer_servings                 float64
spirit_servings                int64
wine_servings                 int64
total_litres_of_pure_alcohol   float64
continent                     object
dtype: object
```

```
In [77]: # read a dataset of Chipotle orders into a DataFrame
orders = pd.read_table('http://bit.ly/chiporders')
orders.head()
```

```
Out[77]:
```

	order_id	quantity	item_name	choice_description	item_price
0	1	1	Chips and Fresh Tomato Salsa	NaN	\$2.39
1	1	1	Izze	[Clementine]	\$3.39
2	1	1	Nantucket Nectar	[Apple]	\$3.39
3	1	1	Chips and Tomatillo-Green Chili Salsa	NaN	\$2.39
4	2	2	Chicken Bowl	[Tomatillo-Red Chili Salsa (Hot), [Black Beans...	\$16.98

```
In [78]: # examine the data type of each Series
orders.dtypes
```

```
Out[78]: order_id                int64
quantity                int64
item_name                object
choice_description       object
item_price               object
dtype: object
```

```
In [79]: # convert a string to a number in order to do math
orders.item_price.str.replace('$', '').astype(float).mean()
```

```
Out[79]: 7.464335785374397
```

```
In [80]: # string method 'contains' checks for a substring and returns a boolean Series
orders.item_name.str.contains('Chicken').head()
```

```
Out[80]: 0    False
1    False
2    False
3    False
4     True
Name: item_name, dtype: bool
```

```
In [81]: # convert a boolean Series to an integer (False = 0, True = 1)  
orders.item_name.str.contains('Chicken').astype(int).head()
```

```
Out[81]: 0    0  
1    0  
2    0  
3    0  
4    1  
Name: item_name, dtype: int32
```

When should I use a "groupby" in pandas?

```
In [82]: # read a dataset of alcohol consumption into a DataFrame  
import pandas as pd  
drinks = pd.read_csv('http://bit.ly/drinksbycountry')  
drinks.head()
```

```
Out[82]:
```

	country	beer_servings	spirit_servings	wine_servings	total_litres_of_pure_alcohol	continent
0	Afghanistan	0	0	0	0.0	
1	Albania	89	132	54	4.9	Eu
2	Algeria	25	0	14	0.7	A
3	Andorra	245	138	312	12.4	Eu
4	Angola	217	57	45	5.9	A

```
In [83]: drinks.describe()
```

```
Out[83]:
```

	beer_servings	spirit_servings	wine_servings	total_litres_of_pure_alcohol
count	193.000000	193.000000	193.000000	193.000000
mean	106.160622	80.994819	49.450777	4.717098
std	101.143103	88.284312	79.697598	3.773298
min	0.000000	0.000000	0.000000	0.000000
25%	20.000000	4.000000	1.000000	1.300000
50%	76.000000	56.000000	8.000000	4.200000
75%	188.000000	128.000000	59.000000	7.200000
max	376.000000	438.000000	370.000000	14.400000

```
In [84]: # calculate the mean beer servings across the entire dataset  
drinks.beer_servings.mean()
```

```
Out[84]: 106.16062176165804
```

```
In [85]: # calculate the mean beer servings just for countries in Africa  
drinks[drinks.continent=='Africa'].beer_servings.mean()
```

```
Out[85]: 61.471698113207545
```

```
In [86]: drinks.continent.unique()
```

```
Out[86]: array(['Asia', 'Europe', 'Africa', 'North America', 'South America',  
               'Oceania'], dtype=object)
```

```
In [87]: # calculate the mean beer servings for each continent  
drinks.groupby('continent').beer_servings.mean()
```

```
Out[87]: continent  
Africa          61.471698  
Asia            37.045455  
Europe          193.777778  
North America   145.434783  
Oceania         89.687500  
South America   175.083333  
Name: beer_servings, dtype: float64
```

Documentation for **groupby** [_\(http://pandas.pydata.org/pandas-docs/stable/generated/pandas.DataFrame.groupby.html\)](http://pandas.pydata.org/pandas-docs/stable/generated/pandas.DataFrame.groupby.html)

```
In [88]: # other aggregation functions (such as 'max') can also be used with groupby  
drinks.groupby('continent').beer_servings.max()
```

```
Out[88]: continent  
Africa          376  
Asia            247  
Europe          361  
North America   285  
Oceania         306  
South America   333  
Name: beer_servings, dtype: int64
```



```
In [89]: # multiple aggregation functions can be applied simultaneously
drinks.groupby('continent').agg(['count', 'mean', 'min', 'max'])
```

Out[89]:

	beer_servings				spirit_servings				wine_servings			
	count	mean	min	max	count	mean	min	max	count	mean	min	max
continent												
Africa	53	61.471698	0	376	53	16.339623	0	152	53	16.264151		
Asia	44	37.045455	0	247	44	60.840909	0	326	44	9.068182		
Europe	45	193.777778	0	361	45	132.555556	0	373	45	142.222222		
North America	23	145.434783	1	285	23	165.739130	68	438	23	24.521739		
Oceania	16	89.687500	0	306	16	58.437500	0	254	16	35.625000		
South America	12	175.083333	93	333	12	114.750000	25	302	12	62.416667		

Documentation for `agg` (<http://pandas.pydata.org/pandas-docs/stable/generated/pandas.core.groupby.DataFrameGroupBy.agg.html>)

```
In [90]: # specifying a column to which the aggregation function should be applied is
drinks.groupby('continent').mean()
```

Out[90]:

	beer_servings	spirit_servings	wine_servings	total_litres_of_pure_alcohol
continent				
Africa	61.471698	16.339623	16.264151	3.007547
Asia	37.045455	60.840909	9.068182	2.170455
Europe	193.777778	132.555556	142.222222	8.617778
North America	145.434783	165.739130	24.521739	5.995652
Oceania	89.687500	58.437500	35.625000	3.381250
South America	175.083333	114.750000	62.416667	6.308333

```
In [91]: drinks.describe()
```

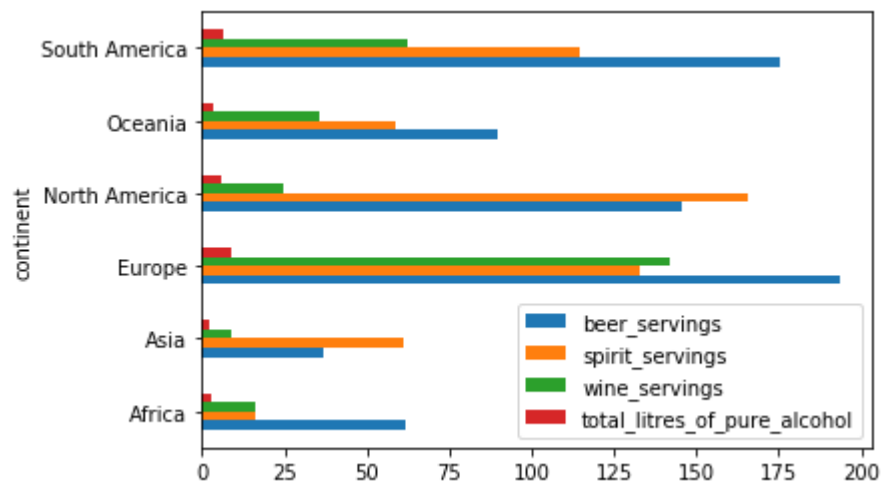
Out[91]:

	beer_servings	spirit_servings	wine_servings	total_litres_of_pure_alcohol
count	193.000000	193.000000	193.000000	193.000000
mean	106.160622	80.994819	49.450777	4.717098
std	101.143103	88.284312	79.697598	3.773298
min	0.000000	0.000000	0.000000	0.000000
25%	20.000000	4.000000	1.000000	1.300000
50%	76.000000	56.000000	8.000000	4.200000
75%	188.000000	128.000000	59.000000	7.200000
max	376.000000	438.000000	370.000000	14.400000

```
In [92]: # allow plots to appear in the notebook
import matplotlib.pyplot as plt
%matplotlib inline
```


```
In [93]: # side-by-side bar plot of the DataFrame directly above
drinks.groupby('continent').mean().plot(kind='barh')
```

Out[93]: <matplotlib.axes._subplots.AxesSubplot at 0x11ff729a3c8>



Documentation for **plot** (<http://pandas.pydata.org/pandas-docs/stable/generated/pandas.DataFrame.plot.html>)

[[Back to top](#)]

```
In [94]:  # count how many times each value in the Series occurs  
movies.genre.value_counts()
```

```
Out[94]: Drama          278  
         Comedy         156  
         Action         136  
         Crime          124  
         Biography       77  
         Adventure       75  
         Animation       62  
         Horror          29  
         Mystery         16  
         Western          9  
         Thriller         5  
         Sci-Fi          5  
         Film-Noir        3  
         Family           2  
         Fantasy          1  
         History          1  
         Name: genre, dtype: int64
```

Documentation for **value_counts** [_\(http://pandas.pydata.org/pandas-docs/stable/generated/pandas.Series.value_counts.html\)](http://pandas.pydata.org/pandas-docs/stable/generated/pandas.Series.value_counts.html)

In [95]: `movies.head(10)`

Out[95]:

	star_rating	title	content_rating	genre	duration	actors_list
0	9.3	The Shawshank Redemption	R	Crime	142	[u'Tim Robbins', u'Morgan Freeman', u'Bob Gunt...
1	9.2	The Godfather	R	Crime	175	[u'Marlon Brando', u'Al Pacino', u'James Caan']
2	9.1	The Godfather: Part II	R	Crime	200	[u'Al Pacino', u'Robert De Niro', u'Robert Duv...
3	9.0	The Dark Knight	PG-13	Action	152	[u'Christian Bale', u'Heath Ledger', u'Aaron E...
4	8.9	Pulp Fiction	R	Crime	154	[u'John Travolta', u'Uma Thurman', u'Samuel L....
5	8.9	12 Angry Men	NOT RATED	Drama	96	[u'Henry Fonda', u'Lee J. Cobb', u'Martin Bals...
6	8.9	The Good, the Bad and the Ugly	NOT RATED	Western	161	[u'Clint Eastwood', u'Eli Wallach', u'Lee Van ...
7	8.9	The Lord of the Rings: The Return of the King	PG-13	Adventure	201	[u'Elijah Wood', u'Viggo Mortensen', u'Ian McK...
8	8.9	Schindler's List	R	Biography	195	[u'Liam Neeson', u'Ralph Fiennes', u'Ben Kings...
9	8.9	Fight Club	R	Drama	139	[u'Brad Pitt', u'Edward Norton', u'Helena Bonh...

```
In [96]: ▶ # display percentages instead of raw counts
movies.genre.value_counts(normalize=True)
```

```
Out[96]: Drama      0.283963
Comedy      0.159346
Action      0.138917
Crime       0.126660
Biography   0.078652
Adventure   0.076609
Animation   0.063330
Horror      0.029622
Mystery     0.016343
Western     0.009193
Thriller    0.005107
Sci-Fi      0.005107
Film-Noir   0.003064
Family      0.002043
Fantasy     0.001021
History     0.001021
Name: genre, dtype: float64
```

```
In [97]: ▶ # 'value_counts' (like many pandas methods) outputs a Series
type(movies.genre.value_counts())
```

```
Out[97]: pandas.core.series.Series
```

```
In [98]: ▶ # display the unique values in the Series
movies.genre.unique()
```

```
Out[98]: array(['Crime', 'Action', 'Drama', 'Western', 'Adventure', 'Biography',
                'Comedy', 'Animation', 'Mystery', 'Horror', 'Film-Noir', 'Sci-Fi',
                'History', 'Thriller', 'Family', 'Fantasy'], dtype=object)
```

```
In [99]: ▶ # count the number of unique values in the Series
movies.genre.nunique()
```

```
Out[99]: 16
```

Documentation for **unique** (<http://pandas.pydata.org/pandas-docs/stable/generated/pandas.Series.unique.html>) and **nunique** (<http://pandas.pydata.org/pandas-docs/stable/generated/pandas.Series.nunique.html>).

```
In [100]: ▶ print(movies.content_rating.nunique())
movies.genre.nunique()
```

```
12
```

```
Out[100]: 16
```

```
In [101]: # compute a cross-tabulation of two Series
pd.crosstab(movies.genre, movies.content_rating, margins=True)
```

Out[101]:

content_rating	APPROVED	G	GP	NC-17	NOT RATED	PASSED	PG	PG-13	R	TV-MA	UNRATED
genre											
Action	3	1	1	0	4	1	11	44	67	0	3
Adventure	3	2	0	0	5	1	21	23	17	0	2
Animation	3	20	0	0	3	0	25	5	5	0	1
Biography	1	2	1	0	1	0	6	29	36	0	0
Comedy	9	2	1	1	16	3	23	23	73	0	4
Crime	6	0	0	1	7	1	6	4	87	0	11
Drama	12	3	0	4	24	1	25	55	143	1	9
Family	0	1	0	0	0	0	1	0	0	0	0
Fantasy	0	0	0	0	0	0	0	0	1	0	0
Film-Noir	1	0	0	0	1	0	0	0	0	0	1
History	0	0	0	0	0	0	0	0	0	0	1
Horror	2	0	0	1	1	0	1	2	16	0	5
Mystery	4	1	0	0	1	0	1	2	6	0	1
Sci-Fi	1	0	0	0	0	0	0	1	3	0	0
Thriller	1	0	0	0	0	0	1	0	3	0	0
Western	1	0	0	0	2	0	2	1	3	0	0
All	47	32	3	7	65	7	123	189	460	1	38

Documentation for **crosstab** (<http://pandas.pydata.org/pandas-docs/stable/generated/pandas.crosstab.html>)

Exploring a numeric Series:

```
In [102]: # calculate various summary statistics
movies.duration.describe()
```

```
Out[102]: count    979.000000
mean    120.979571
std     26.218010
min     64.000000
25%    102.000000
50%    117.000000
75%    134.000000
max    242.000000
Name: duration, dtype: float64
```

```
In [103]: ▶ # many statistics are implemented as Series methods
          movies.duration.mean()
```

```
Out[103]: 120.97957099080695
```

Documentation for `mean` (<http://pandas.pydata.org/pandas-docs/stable/generated/pandas.Series.mean.html>)


```
In [104]: ▶ # 'value_counts' is primarily useful for categorical data, not numerical data
          movies.duration.value_counts().head()
```

```
Out[104]: 112    23
          113    22
          102    20
          101    20
          129    19
          Name: duration, dtype: int64
```

```
In [105]: ▶ # allow plots to appear in the notebook
          %matplotlib inline
```


```
In [106]: ▶ # histogram of the 'duration' Series (shows the distribution of a numerical variable)
          movies.duration.plot(kind='hist')
```

```
Out[106]: <matplotlib.axes._subplots.AxesSubplot at 0x11ff8beabc8>
```



```
In [107]: # bar plot of the 'value_counts' for the 'genre' Series  
movies.genre.value_counts().plot(kind='bar')
```

```
Out[107]: <matplotlib.axes._subplots.AxesSubplot at 0x11ff8c87488>
```



Documentation for **plot** (<http://pandas.pydata.org/pandas-docs/stable/generated/pandas.Series.plot.html>).

How do I handle missing values in pandas?


```
In [108]: # read a dataset of UFO reports into a DataFrame
ufo = pd.read_csv('http://bit.ly/uforeports')
ufo.tail()
```

Out[108]:

	City	Colors Reported	Shape Reported	State	Time
18236	Grant Park	NaN	TRIANGLE	IL	12/31/2000 23:00
18237	Spirit Lake	NaN	DISK	IA	12/31/2000 23:00
18238	Eagle River	NaN	NaN	WI	12/31/2000 23:45
18239	Eagle River	RED	LIGHT	WI	12/31/2000 23:45
18240	Ybor	NaN	OVAL	FL	12/31/2000 23:59

What does "NaN" mean?

- "NaN" is not a string, rather it's a special value: `numpy.nan`.
- It stands for "Not a Number" and indicates a **missing value**.
- `read_csv` detects missing values (by default) when reading the file, and replaces them with this special value.

Documentation for `read_csv` (http://pandas.pydata.org/pandas-docs/stable/generated/pandas.read_csv.html)

```
In [109]: # 'isnull' returns a DataFrame of booleans (True if missing, False if not missing)
ufo.isnull().tail()
```

Out[109]:

	City	Colors Reported	Shape Reported	State	Time
18236	False	True	False	False	False
18237	False	True	False	False	False
18238	False	True	True	False	False
18239	False	False	False	False	False
18240	False	True	False	False	False

```
In [110]: # 'nonnull' returns the opposite of 'isnull' (True if not missing, False if missing)
ufo.nonnull().tail()
```

Out[110]:

	City	Colors Reported	Shape Reported	State	Time
18236	True	False	True	True	True
18237	True	False	True	True	True
18238	True	False	False	True	True
18239	True	True	True	True	True
18240	True	False	True	True	True

Documentation for `isnull` (<http://pandas.pydata.org/pandas-docs/stable/generated/pandas.isnull.html>)

[docs/stable/generated/pandas.DataFrame.isnull.html](https://pandas.pydata.org/pandas-docs/stable/generated/pandas.DataFrame.isnull.html)) and **notnull** (<http://pandas.pydata.org/pandas-docs/stable/generated/pandas.notnull.html>)

```
In [111]: # count the number of missing values in each Series  
ufo.isnull().sum()
```

```
Out[111]: City                25  
Colors Reported            15359  
Shape Reported            2644  
State                      0  
Time                      0  
dtype: int64
```

This calculation works because:

1. The **sum** method for a DataFrame operates on **axis=0** by default (and thus produces column sums).
2. In order to add boolean values, pandas converts **True** to **1** and **False** to **0**.

```
In [112]: # use the 'isnull' Series method to filter the DataFrame rows  
ufo[ufo.City.isnull()].head()
```

```
Out[112]:
```

	City	Colors Reported	Shape Reported	State	Time
21	NaN	NaN	NaN	LA	8/15/1943 0:00
22	NaN	NaN	LIGHT	LA	8/15/1943 0:00
204	NaN	NaN	DISK	CA	7/15/1952 12:30
241	NaN	BLUE	DISK	MT	7/4/1953 14:00
613	NaN	NaN	DISK	NV	7/1/1960 12:00

How to handle missing values depends on the dataset as well as the nature of your analysis. Here are some options:

```
In [113]: # examine the number of rows and columns  
ufo.shape
```

```
Out[113]: (18241, 5)
```

```
In [114]: # if 'any' values are missing in a row, then drop that row  
ufo.dropna(how='any').shape
```

```
Out[114]: (2486, 5)
```

Documentation for **dropna** (<http://pandas.pydata.org/pandas-docs/stable/generated/pandas.DataFrame.dropna.html>)

```
In [115]: ▶ # 'inplace' parameter for 'dropna' is False by default, thus rows were only c
ufo.shape
```

```
Out[115]: (18241, 5)
```

```
In [116]: ▶ # if 'all' values are missing in a row, then drop that row (none are dropped)
ufo.dropna(how='all').shape
```

```
Out[116]: (18241, 5)
```

```
In [117]: ▶ # if 'any' values are missing in a row (considering only 'City' and 'Shape Re
ufo.dropna(subset=['City', 'Shape Reported'], how='any').shape
```

```
Out[117]: (15576, 5)
```

```
In [118]: ▶ # if 'all' values are missing in a row (considering only 'City' and 'Shape Re
ufo.dropna(subset=['City', 'Shape Reported'], how='all').shape
```

```
Out[118]: (18237, 5)
```

```
In [119]: ▶ # 'value_counts' does not include missing values by default
ufo['Shape Reported'].value_counts().head()
```

```
Out[119]: LIGHT      2803
          DISK       2122
          TRIANGLE   1889
          OTHER      1402
          CIRCLE     1365
          Name: Shape Reported, dtype: int64
```

```
In [120]: ▶ # explicitly include missing values
ufo['City'].value_counts(dropna=True).head()
```

```
Out[120]: Seattle      187
          New York City  161
          Phoenix      137
          Houston      108
          Las Vegas     105
          Name: City, dtype: int64
```

Documentation for `value_counts` (http://pandas.pydata.org/pandas-docs/stable/generated/pandas.Series.value_counts.html)

```
In [121]: ▶ # fill in missing values with a specified value
ufo['Shape Reported'].fillna(value='VARIOUS', inplace=True)
```

Documentation for `fillna` (<http://pandas.pydata.org/pandas-docs/stable/generated/pandas.DataFrame.fillna.html>)

```
In [122]: ▶ # confirm that the missing values were filled in
          ufo['Shape Reported'].value_counts().head()
```

```
Out[122]: VARIOUS      2977
          LIGHT       2803
          DISK        2122
          TRIANGLE    1889
          OTHER       1402
          Name: Shape Reported, dtype: int64
```

[Working with missing data in pandas \(http://pandas.pydata.org/pandas-docs/stable/missing_data.html\)](http://pandas.pydata.org/pandas-docs/stable/missing_data.html)

What do I need to know about the pandas index?

```
In [123]: ▶ # read a dataset of alcohol consumption into a DataFrame
          drinks = pd.read_csv('http://bit.ly/drinksbycountry')
          drinks.head()
```

```
Out[123]:
```

	country	beer_servings	spirit_servings	wine_servings	total_litres_of_pure_alcohol	continent
0	Afghanistan	0	0	0	0.0	
1	Albania	89	132	54	4.9	Eu
2	Algeria	25	0	14	0.7	A
3	Andorra	245	138	312	12.4	Eu
4	Angola	217	57	45	5.9	A

```
In [124]: ▶ # every DataFrame has an index (sometimes called the "row labels")
          drinks.index
```


```
Out[124]: RangeIndex(start=0, stop=193, step=1)
```

```
In [125]: ▶ # column names are also stored in a special "index" object
          drinks.columns
```

```
Out[125]: Index(['country', 'beer_servings', 'spirit_servings', 'wine_servings',
                  'total_litres_of_pure_alcohol', 'continent'],
                  dtype='object')
```

```
In [126]: ▶ # neither the index nor the columns are included in the shape
          drinks.shape
```

```
Out[126]: (193, 6)
```

```
In [127]:  # index and columns both default to integers if you don't define them
pd.read_table('http://bit.ly/movieusers', header=None, sep='|').head()
```

Out[127]:

	0	1	2		3	4
0	1	24	M	technician	85711	
1	2	53	F	other	94043	
2	3	23	M	writer	32067	
3	4	24	M	technician	43537	
4	5	33	F	other	15213	

What is the index used for?

1. identification
2. selection
3. alignment

```
In [128]: # identification: index remains with each row when filtering the DataFrame  
drinks[drinks.continent=='South America']
```

Out[128]:

	country	beer_servings	spirit_servings	wine_servings	total_litres_of_pure_alcohol	continent
6	Argentina	193	25	221	8.3	South America
20	Bolivia	167	41	8	3.8	South America
23	Brazil	245	145	16	7.2	South America
35	Chile	130	124	172	7.6	South America
37	Colombia	159	76	3	4.2	South America
52	Ecuador	162	74	3	4.2	South America
72	Guyana	93	302	1	7.1	South America
132	Paraguay	213	117	74	7.3	South America
133	Peru	163	160	21	6.1	South America
163	Suriname	128	178	7	5.6	South America
185	Uruguay	115	35	220	6.6	South America
188	Venezuela	333	100	3	7.7	South America

```
In [129]: # selection: select a portion of the DataFrame using the index  
drinks.loc[0:10, 'beer_servings']
```

Out[129]:

0	0
1	89
2	25
3	245
4	217
5	102
6	193
7	21
8	261
9	279
10	21

Name: beer_servings, dtype: int64

Documentation for `loc` [_\(http://pandas.pydata.org/pandas-docs/stable/generated/pandas.DataFrame.loc.html\)](http://pandas.pydata.org/pandas-docs/stable/generated/pandas.DataFrame.loc.html)

```
In [130]: # set an existing column as the index
drinks.set_index('country', inplace=True)
drinks.head()
```

Out[130]:

	beer_servings	spirit_servings	wine_servings	total_litres_of_pure_alcohol	continent
country					
Afghanistan	0	0	0	0.0	Asia
Albania	89	132	54	4.9	Europe
Algeria	25	0	14	0.7	Africa
Andorra	245	138	312	12.4	Europe
Angola	217	57	45	5.9	Africa

Documentation for `set_index` [_\(http://pandas.pydata.org/pandas-docs/stable/generated/pandas.DataFrame.set_index.html\)](http://pandas.pydata.org/pandas-docs/stable/generated/pandas.DataFrame.set_index.html)

```
In [131]: # 'country' is now the index
drinks.index
```

Out[131]: Index(['Afghanistan', 'Albania', 'Algeria', 'Andorra', 'Angola',
'Antigua & Barbuda', 'Argentina', 'Armenia', 'Australia', 'Austria',
...,
'Tanzania', 'USA', 'Uruguay', 'Uzbekistan', 'Vanuatu', 'Venezuela',
'Vietnam', 'Yemen', 'Zambia', 'Zimbabwe'],
dtype='object', name='country', length=193)

```
In [132]: # 'country' is no longer a column
drinks.columns
```

Out[132]: Index(['beer_servings', 'spirit_servings', 'wine_servings',
'total_litres_of_pure_alcohol', 'continent'],
dtype='object')

```
In [133]: # 'country' data is no longer part of the DataFrame contents
drinks.shape
```

Out[133]: (193, 5)

```
In [134]: # country name can now be used for selection
drinks.loc['Brazil', 'beer_servings']
```

Out[134]: 245

```
In [135]: # index name is optional
drinks.index.name = None
drinks.head()
```

Out[135]:

	beer_servings	spirit_servings	wine_servings	total_litres_of_pure_alcohol	continent
Afghanistan	0	0	0	0.0	Asia
Albania	89	132	54	4.9	Europe
Algeria	25	0	14	0.7	Africa
Andorra	245	138	312	12.4	Europe
Angola	217	57	45	5.9	Africa

```
In [136]: # restore the index name, and move the index back to a column
drinks.index.name = 'country'
drinks.reset_index(inplace=True)
drinks.head()
```

Out[136]:

	country	beer_servings	spirit_servings	wine_servings	total_litres_of_pure_alcohol	continent
0	Afghanistan	0	0	0	0.0	Asia
1	Albania	89	132	54	4.9	Europe
2	Algeria	25	0	14	0.7	Africa
3	Andorra	245	138	312	12.4	Europe
4	Angola	217	57	45	5.9	Africa

Documentation for **reset_index** (http://pandas.pydata.org/pandas-docs/stable/generated/pandas.DataFrame.reset_index.html)

```
In [137]: # many DataFrame methods output a DataFrame
drinks.describe()
```

Out[137]:

	beer_servings	spirit_servings	wine_servings	total_litres_of_pure_alcohol
count	193.000000	193.000000	193.000000	193.000000
mean	106.160622	80.994819	49.450777	4.717098
std	101.143103	88.284312	79.697598	3.773298
min	0.000000	0.000000	0.000000	0.000000
25%	20.000000	4.000000	1.000000	1.300000
50%	76.000000	56.000000	8.000000	4.200000
75%	188.000000	128.000000	59.000000	7.200000
max	376.000000	438.000000	370.000000	14.400000


```
In [138]: ▶ # you can interact with any DataFrame using its index and columns
drinks.describe().loc['25%', 'beer_servings']
```

Out[138]: 20.0

[Indexing and selecting data \(http://pandas.pydata.org/pandas-docs/stable/indexing.html\)](http://pandas.pydata.org/pandas-docs/stable/indexing.html)

[\[Back to top\]](#)

How do I select multiple rows and columns from a pandas DataFrame?

```
In [139]: ▶ # read a dataset of UFO reports into a DataFrame
ufo = pd.read_csv('http://bit.ly/uforeports')
ufo.head(3)
```

Out[139]:

	City	Colors Reported	Shape Reported	State	Time
0	Ithaca	NaN	TRIANGLE	NY	6/1/1930 22:00
1	Willingboro	NaN	OTHER	NJ	6/30/1930 20:00
2	Holyoke	NaN	OVAl	CO	2/15/1931 14:00

The `loc` (<http://pandas.pydata.org/pandas-docs/stable/generated/pandas.DataFrame.loc.html>) method is used to select rows and columns by **label**. You can pass it:

- A single label
- A list of labels
- A slice of labels
- A boolean Series
- A colon (which indicates "all labels")

```
In [140]: ▶ # row 0, all columns
ufo.loc[0, :]
```

Out[140]:

City	Ithaca
Colors Reported	NaN
Shape Reported	TRIANGLE
State	NY
Time	6/1/1930 22:00
Name: 0, dtype: object	

```
In [141]: # rows 0 and 1 and 2, all columns
ufo.loc[[0, 1, 2], :]
```

Out[141]:

	City	Colors Reported	Shape Reported	State	Time
0	Ithaca	NaN	TRIANGLE	NY	6/1/1930 22:00
1	Willingboro	NaN	OTHER	NJ	6/30/1930 20:00
2	Holyoke	NaN	OVAL	CO	2/15/1931 14:00

```
In [142]: # rows 0 through 2 (inclusive), all columns
ufo.loc[0:2, :]
```

Out[142]:

	City	Colors Reported	Shape Reported	State	Time
0	Ithaca	NaN	TRIANGLE	NY	6/1/1930 22:00
1	Willingboro	NaN	OTHER	NJ	6/30/1930 20:00
2	Holyoke	NaN	OVAL	CO	2/15/1931 14:00

```
In [143]: # this implies "all columns", but explicitly stating "all columns" is better
ufo.loc[0:2]
```

Out[143]:

	City	Colors Reported	Shape Reported	State	Time
0	Ithaca	NaN	TRIANGLE	NY	6/1/1930 22:00
1	Willingboro	NaN	OTHER	NJ	6/30/1930 20:00
2	Holyoke	NaN	OVAL	CO	2/15/1931 14:00

```
In [144]: # rows 0 through 2 (inclusive), column 'City'
ufo.loc[0:2, 'City']
```

Out[144]:

0	Ithaca
1	Willingboro
2	Holyoke

Name: City, dtype: object

```
In [145]: # rows 0 through 2 (inclusive), columns 'City' and 'State'
ufo.loc[0:2, ['City', 'State']]
```

Out[145]:

	City	State
0	Ithaca	NY
1	Willingboro	NJ
2	Holyoke	CO

```
In [146]: # accomplish the same thing using double brackets - but using 'loc' is preferred  
ufo[['City', 'State']].head(3)
```

Out[146]:

	City	State
0	Ithaca	NY
1	Willingboro	NJ
2	Holyoke	CO

```
In [147]: # rows 0 through 2 (inclusive), columns 'City' through 'State' (inclusive)  
ufo.loc[0:2, 'City':'State']
```

Out[147]:

	City	Colors Reported	Shape Reported	State
0	Ithaca	NaN	TRIANGLE	NY
1	Willingboro	NaN	OTHER	NJ
2	Holyoke	NaN	OVAL	CO

```
In [148]: # accomplish the same thing using 'head' and 'drop'  
ufo.head(3).drop('Time', axis=1)
```

Out[148]:

	City	Colors Reported	Shape Reported	State
0	Ithaca	NaN	TRIANGLE	NY
1	Willingboro	NaN	OTHER	NJ
2	Holyoke	NaN	OVAL	CO

```
In [149]: # rows in which the 'City' is 'Oakland', column 'State'  
ufo.loc[ufo.City=='Oakland', 'State']
```

Out[149]:

1694	CA
2144	CA
4686	MD
7293	CA
8488	CA
8768	CA
10816	OR
10948	CA
11045	CA
12322	CA
12941	CA
16803	MD
17322	CA

Name: State, dtype: object

```
In [150]: # accomplish the same thing using "chained indexing" - but using 'loc' is preferred
ufo[ufo.City=='Oakland'].State
```

```
Out[150]: 1694      CA
          2144      CA
          4686      MD
          7293      CA
          8488      CA
          8768      CA
          10816     OR
          10948      CA
          11045      CA
          12322      CA
          12941      CA
          16803      MD
          17322      CA
          Name: State, dtype: object
```

The `[**`iloc`**]` (<http://pandas.pydata.org/pandas-docs/stable/generated/pandas.DataFrame.iloc.html>) method is used to select rows and columns by **integer position**. You can pass it:

- A single integer position
- A list of integer positions
- A slice of integer positions
- A colon (which indicates "all integer positions")

```
In [151]: ufo.head()
```

```
Out[151]:
```

	City	Colors Reported	Shape Reported	State	Time
0	Ithaca	NaN	TRIANGLE	NY	6/1/1930 22:00
1	Willingboro	NaN	OTHER	NJ	6/30/1930 20:00
2	Holyoke	NaN	OVAL	CO	2/15/1931 14:00
3	Abilene	NaN	DISK	KS	6/1/1931 13:00
4	New York Worlds Fair	NaN	LIGHT	NY	4/18/1933 19:00

```
In [152]: # rows in positions 0 and 1, columns in positions 0 and 3
ufo.iloc[[0, 1], [0, 3]]
```

```
Out[152]:
```

	City	State
0	Ithaca	NY
1	Willingboro	NJ

```
In [153]: ▶ # rows in positions 0 through 2 (exclusive), columns in positions 0 through 4
ufo.iloc[0:2, 0:4]
```

Out[153]:

	City	Colors Reported	Shape Reported	State
0	Ithaca	NaN	TRIANGLE	NY
1	Willingboro	NaN	OTHER	NJ

```
In [154]: ▶ # rows in positions 0 through 2 (exclusive), all columns
ufo.iloc[0:2, :]
```

Out[154]:

	City	Colors Reported	Shape Reported	State	Time
0	Ithaca	NaN	TRIANGLE	NY	6/1/1930 22:00
1	Willingboro	NaN	OTHER	NJ	6/30/1930 20:00

```
In [155]: ▶ # accomplish the same thing - but using 'iloc' is preferred since it's more explicit
ufo[0:2]
```

Out[155]:

	City	Colors Reported	Shape Reported	State	Time
0	Ithaca	NaN	TRIANGLE	NY	6/1/1930 22:00
1	Willingboro	NaN	OTHER	NJ	6/30/1930 20:00

The `[**`ix`**]` (<http://pandas.pydata.org/pandas-docs/stable/generated/pandas.DataFrame.ix.html>) method is used to select rows and columns by **label or integer position**, and should only be used when you need to mix label-based and integer-based selection in the same call.

```
In [305]: ▶ # read a dataset of alcohol consumption into a DataFrame and set 'country' as index
import pandas as pd
drinks = pd.read_csv('http://bit.ly/drinksbycountry', index_col='country')
drinks.head()
```

Out[305]:

	beer_servings	spirit_servings	wine_servings	total_litres_of_pure_alcohol	continent
country					
Afghanistan	0	0	0	0.0	Asia
Albania	89	132	54	4.9	Europe
Algeria	25	0	14	0.7	Africa
Andorra	245	138	312	12.4	Europe
Angola	217	57	45	5.9	Africa

```
In [282]: ▶ # row with label 'Albania', column in position 0
#drinks.ix['Andorra', 0]--->Since IX is now not in use
```

```
In [284]: ▶ # row in position 1, column with label 'beer_servings'
drinks.ix[1, 'beer_servings']
```

C:\Users\Meghaa\Anaconda3\lib\site-packages\ipykernel_launcher.py:2: Future Warning:
.ix is deprecated. Please use
.loc for label based indexing or
.iloc for positional indexing

See the documentation here:
http://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#ix-indexer-is-deprecated (http://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#ix-indexer-is-deprecated)

Out[284]: 89.0

****Rules for using numbers with `ix`:****

- If the index is ****strings****, numbers are treated as ****integer positions****, and thus slices are ****exclusive**** on the right.
- If the index is ****integers****, numbers are treated as ****labels****, and thus slices are ****inclusive****.

```
In [298]: ▶ # rows 'Albania' through 'Andorra' (inclusive), columns in positions 0 through
drinks.ix['Albania':'Andorra', 0:2]
```

C:\Users\Meghaa\Anaconda3\lib\site-packages\ipykernel_launcher.py:2: Future Warning:
.ix is deprecated. Please use
.loc for label based indexing or
.iloc for positional indexing

See the documentation here:
http://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#ix-indexer-is-deprecated (http://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#ix-indexer-is-deprecated)

Out[298]:

<u>country</u>	<u>beer_servings</u>
----------------	----------------------

```
In [299]: # rows 0 through 2 (inclusive), columns in positions 0 through 2 (exclusive)  
ufo.ix[0:2, 0:2]
```

C:\Users\Meghaa\Anaconda3\lib\site-packages\ipykernel_launcher.py:2: Future Warning:

.ix is deprecated. Please use
.loc for label based indexing or
.iloc for positional indexing

See the documentation here:

http://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#ix-indexer-is-deprecated (http://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#ix-indexer-is-deprecated)

Out[299]:

	City	Colors Reported
0	Ithaca	NaN
1	Willingboro	NaN
2	Holyoke	NaN

```
In [290]: # use the 'category' data type (new in pandas 0.15) to store the 'continent'  
drinks['continent'] = drinks.continent.astype('category')  
drinks.dtypes
```

Out[290]:

country	object
beer_servings	float64
spirit_servings	float64
wine_servings	float64
total_litres_of_pure_alcohol	float64
continent	category
dtype:	object

```
In [291]: # 'continent' Series appears to be unchanged  
drinks.continent.head()
```

Out[291]:

0	Asia
1	Europe
2	Africa
3	Europe
4	Africa

Name: continent, dtype: category
Categories (6, object): [Africa, Asia, Europe, North America, Oceania, South America]

```
In [163]: ▶ # strings are now encoded (0 means 'Africa', 1 means 'Asia', 2 means 'Europe')
drinks.continent.cat.codes.head()
```

```
Out[163]: country
Afghanistan    1
Albania        2
Algeria        0
Andorra        2
Angola         0
dtype: int8
```

```
In [164]: ▶ # memory usage has been drastically reduced
drinks.memory_usage(deep=True)
```

```
Out[164]: Index                17708
beer_servings              1544
spirit_servings            1544
wine_servings              1544
total_litres_of_pure_alcohol 1544
continent                  744
dtype: int64
```

```
In [315]: ▶ # repeat this process for the 'country' Series
drinks['country'] = drinks.country.astype('category')
drinks.memory_usage(deep=True)
```

```
-----
AttributeError                                Traceback (most recent call last)
<ipython-input-315-6edf4187ca71> in <module>
      1 # repeat this process for the 'country' Series
----> 2 drinks['country'] = drinks.country.astype('category')
      3 drinks.memory_usage(deep=True)

~\Anaconda3\lib\site-packages\pandas\core\generic.py in __getattr__(self, name)
    5177         if self._info_axis._can_hold_identifiers_and_holds_name
    (name):
    5178             return self[name]
-> 5179         return object.__getattribute__(self, name)
    5180
    5181     def __setattr__(self, name, value):

AttributeError: 'DataFrame' object has no attribute 'country'
```



```
In [304]: # memory usage increased because we created 193 categories
drinks.country.cat.categories
```

```
-----
AttributeError                                Traceback (most recent call last)
<ipython-input-304-1af6b0e83ad5> in <module>
      1 # memory usage increased because we created 193 categories
----> 2 drinks.country.cat.categories

~\Anaconda3\lib\site-packages\pandas\core\generic.py in __getattr__(self, name)
    5177         if self._info_axis._can_hold_identifiers_and_holds_name
    (name):
    5178             return self[name]
-> 5179         return object.__getattribute__(self, name)
    5180
    5181     def __setattr__(self, name, value):

AttributeError: 'DataFrame' object has no attribute 'country'
```

The **category** data type should only be used with a string Series that has a **small number of possible values**.

```
In [300]: # create a small DataFrame from a dictionary
df = pd.DataFrame({'ID':[100, 101, 102, 103], 'quality':['good', 'very good',
df
```

Out[300]:

	ID	quality
0	100	good
1	101	very good
2	102	good
3	103	excellent

```
In [301]: # sort the DataFrame by the 'quality' Series (alphabetical order)
df.sort_values('quality')
```

Out[301]:

	ID	quality
3	103	excellent
0	100	good
2	102	good
1	101	very good

```

In [317]: # define a logical ordering for the categories
df['quality'] = df.quality.astype('category', categories=['good', 'very good',
df.quality

-----
--
ValueError                                Traceback (most recent call last)
<ipython-input-317-9cdaae1356ea> in <module>
      1 # define a logical ordering for the categories
----> 2 df['quality'] = df.quality.astype('category', categories=['good',
'very good', 'excellent'], ordered=True)
      3 df.quality

~\Anaconda3\lib\site-packages\pandas\core\generic.py in astype(self, dtype, copy, errors, **kwargs)
    5880         # else, only a single dtype is given
    5881         new_data = self._data.astype(
--> 5882             dtype=dtype, copy=copy, errors=errors, **kwargs
    5883         )
    5884         return self._constructor(new_data).__finalize__(self)

~\Anaconda3\lib\site-packages\pandas\core\internals\managers.py in astype(self, dtype, **kwargs)
    579
    580     def astype(self, dtype, **kwargs):
--> 581         return self.apply("astype", dtype=dtype, **kwargs)
    582
    583     def convert(self, **kwargs):

~\Anaconda3\lib\site-packages\pandas\core\internals\managers.py in apply(self, f, axes, filter, do_integrity_check, consolidate, **kwargs)
    436         kwargs[k] = obj.reindex(b_items, axis=axis, copy=align_copy)
    437
--> 438         applied = getattr(b, f)(**kwargs)
    439         result_blocks = _extend_blocks(applied, result_blocks)
    440

~\Anaconda3\lib\site-packages\pandas\core\internals\blocks.py in astype(self, dtype, copy, errors, values, **kwargs)
    557
    558     def astype(self, dtype, copy=False, errors="raise", values=None, **kwargs):
--> 559         return self._astype(dtype, copy=copy, errors=errors, values=values, **kwargs)
    560
    561     def _astype(self, dtype, copy=False, errors="raise", values=None, **kwargs):

~\Anaconda3\lib\site-packages\pandas\core\internals\blocks.py in _astype(self, dtype, copy, errors, values, **kwargs)
    598         if deprecated_arg in kwargs:
    599             raise ValueError(
--> 600                 "Got an unexpected argument: {}".format(d

```

```
eprecated_arg)  
601  
602 )
```

ValueError: Got an unexpected argument: categories

```
In [310]: ▶ # sort the DataFrame by the 'quality' Series (logical order)  
df.sort_values('quality')
```

Out[310]:

	ID	quality
3	103	excellent
0	100	good
2	102	good
1	101	very good

```
In [316]: ▶ # comparison operators work with ordered categories  
df.loc[df.quality > 'good', :]
```

Out[316]:

	ID	quality
1	101	very good

[Overview of categorical data in pandas \(http://pandas.pydata.org/pandas-docs/stable/categorical.html\)](http://pandas.pydata.org/pandas-docs/stable/categorical.html)

[API reference for categorical methods \(http://pandas.pydata.org/pandas-docs/stable/api.html#categorical\)](http://pandas.pydata.org/pandas-docs/stable/api.html#categorical)

How do I create dummy variables in pandas?

```
In [170]:  # read the training dataset from Kaggle's Titanic competition
train = pd.read_csv('http://bit.ly/kaggletrain')
train.head()
```

Out[170]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	71.2833
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500

```
In [171]: # create the 'Sex_male' dummy variable using the 'map' method
train['Sex_male'] = train.Sex.map({'female':0, 'male':1})
train.head()
```

Out[171]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	71.2833
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500

Documentation for `map` (<http://pandas.pydata.org/pandas-docs/stable/generated/pandas.Series.map.html>)

```
In [172]: # alternative: use 'get_dummies' to create one column for every possible value
pd.get_dummies(train.Sex, prefix = 'Sex', prefix_sep='_', drop_first=True).head()
```

Out[172]:

	Sex_male
0	1
1	0
2	0
3	0
4	1

Generally speaking:

- If you have "**K**" **possible values** for a categorical feature, you only need "**K-1**" **dummy variables** to capture all of the information about that feature.
- One convention is to **drop the first dummy variable**, which defines that level as the "baseline".

```
In [173]: ▶ # drop the first dummy variable ('female') using the 'iloc' method
pd.get_dummies(train.Sex).iloc[:, 1:].head()
```

Out[173]:

	male
0	1
1	0
2	0
3	0
4	1

```
In [174]: ▶ # add a prefix to identify the source of the dummy variables
pd.get_dummies(train.Sex, prefix='Sex').iloc[:, 1:].head()
```

Out[174]:

	Sex_male
0	1
1	0
2	0
3	0
4	1

```
In [175]: ▶ train.Embarked.unique()
```

Out[175]: array(['S', 'C', 'Q', nan], dtype=object)

```
In [176]: # use 'get_dummies' with a feature that has 3 possible values  
pd.get_dummies(train.Embarked, prefix='Embarked').head(10)
```

Out[176]:

	Embarked_C	Embarked_Q	Embarked_S
0	0	0	1
1	1	0	0
2	0	0	1
3	0	0	1
4	0	0	1
5	0	1	0
6	0	0	1
7	0	0	1
8	0	0	1
9	1	0	0

```
In [177]: # drop the first dummy variable ('C')  
pd.get_dummies(train.Embarked, prefix='Embarked').iloc[:, 1:].head(10)
```

Out[177]:

	Embarked_Q	Embarked_S
0	0	1
1	0	0
2	0	1
3	0	1
4	0	1
5	1	0
6	0	1
7	0	1
8	0	1
9	0	0

How to translate these values back to the original 'Embarked' value:


- **0, 0** means **C**
- **1, 0** means **Q**
- **0, 1** means **S**

```
In [178]: # save the DataFrame of dummy variables and concatenate them to the original
embarked_dummies = pd.get_dummies(train.Embarked, prefix='Embarked').iloc[:,
train = pd.concat([train, embarked_dummies], axis=1)
train.head()
```

Out[178]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	71.2833
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500

Documentation for **concat** (<http://pandas.pydata.org/pandas-docs/stable/generated/pandas.concat.html>)


```
In [179]:  # reset the DataFrame
train = pd.read_csv('http://bit.ly/kaggletrain')
train.head()
```

Out[179]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	71.2833
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500

```
In [180]: # pass the DataFrame to 'get_dummies' and specify which columns to dummy (it
pd.get_dummies(train, columns=['Sex', 'Embarked']).head()
```

Out[180]:

	PassengerId	Survived	Pclass	Name	Age	SibSp	Parch	Ticket	Fare	Cabin
0	1	0	3	Braund, Mr. Owen Harris	22.0	1	0	A/5 21171	7.2500	NaN
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	38.0	1	0	PC 17599	71.2833	C85
2	3	1	3	Heikkinen, Miss. Laina	26.0	0	0	STON/O2. 3101282	7.9250	NaN
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	35.0	1	0	113803	53.1000	C123
4	5	0	3	Allen, Mr. William Henry	35.0	0	0	373450	8.0500	NaN

```
In [181]: ▶ from sklearn import preprocessing
le = preprocessing.LabelEncoder()
movies['title'] = le.fit_transform(movies['title'])
movies
```

Out[181]:

	star_rating	title	content_rating	genre	duration	actors_list
0	9.3	866	R	Crime	142	[u'Tim Robbins', u'Morgan Freeman', u'Bob Gunt...
1	9.2	756	R	Crime	175	[u'Marlon Brando', u'Al Pacino', u'James Caan']
2	9.1	757	R	Crime	200	[u'Al Pacino', u'Robert De Niro', u'Robert Duv...
3	9.0	730	PG-13	Action	152	[u'Christian Bale', u'Heath Ledger', u'Aaron E...
4	8.9	560	R	Crime	154	[u'John Travolta', u'Uma Thurman', u'Samuel L....
...
974	7.4	914	PG	Comedy	116	[u'Dustin Hoffman', u'Jessica Lange', u'Teri G...
975	7.4	81	PG	Adventure	118	[u'Michael J. Fox', u'Christopher Lloyd', u'Ma...
976	7.4	459	PG-13	Action	138	[u'Russell Crowe', u'Paul Bettany', u'Billy Bo...
977	7.4	550	PG	Horror	114	[u'JoBeth Williams', u'Heather O'Rourke', u'Cr...
978	7.4	939	R	Crime	126	[u'Charlie Sheen', u'Michael Douglas', u'Tamar...

979 rows × 6 columns

```
In [182]: # use the 'drop_first' parameter (new in pandas 0.18) to drop the first dummy
pd.get_dummies(train, columns=['Sex', 'Embarked'], drop_first=True).head()
```

Out[182]:

	PassengerId	Survived	Pclass	Name	Age	SibSp	Parch	Ticket	Fare	Cabin
0	1	0	3	Braund, Mr. Owen Harris	22.0	1	0	A/5 21171	7.2500	NaN
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	38.0	1	0	PC 17599	71.2833	C85
2	3	1	3	Heikkinen, Miss. Laina	26.0	0	0	STON/O2. 3101282	7.9250	NaN
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	35.0	1	0	113803	53.1000	C123
4	5	0	3	Allen, Mr. William Henry	35.0	0	0	373450	8.0500	NaN

Documentation for `get_dummies` (http://pandas.pydata.org/pandas-docs/stable/generated/pandas.get_dummies.html)

How do I work with dates and times in pandas?

```
In [183]: # read a dataset of UFO reports into a DataFrame
ufo = pd.read_csv('http://bit.ly/uforeports')
ufo.head()
```

Out[183]:

	City	Colors Reported	Shape Reported	State	Time
0	Ithaca	NaN	TRIANGLE	NY	6/1/1930 22:00
1	Willingboro	NaN	OTHER	NJ	6/30/1930 20:00
2	Holyoke	NaN	OVAL	CO	2/15/1931 14:00
3	Abilene	NaN	DISK	KS	6/1/1931 13:00
4	New York Worlds Fair	NaN	LIGHT	NY	4/18/1933 19:00

```
In [184]: ▶ # 'Time' is currently stored as a string
ufo.dtypes
```

```
Out[184]: City                object
Colors Reported             object
Shape Reported              object
State                       object
Time                        object
dtype: object
```

```
In [185]: ▶ # hour could be accessed using string slicing, but this approach breaks too e
ufo.Time.str.slice(-5, -3).astype(int).head()
```

```
Out[185]: 0    22
          1    20
          2    14
          3    13
          4    19
          Name: Time, dtype: int32
```

```
In [186]: ▶ # convert 'Time' to datetime format
ufo['Time'] = pd.to_datetime(ufo.Time)
ufo.head()
```

```
Out[186]:
```

	City	Colors Reported	Shape Reported	State	Time
0	Ithaca	NaN	TRIANGLE	NY	1930-06-01 22:00:00
1	Willingboro	NaN	OTHER	NJ	1930-06-30 20:00:00
2	Holyoke	NaN	OVAL	CO	1931-02-15 14:00:00
3	Abilene	NaN	DISK	KS	1931-06-01 13:00:00
4	New York Worlds Fair	NaN	LIGHT	NY	1933-04-18 19:00:00

```
In [187]: ▶ ufo.dtypes
```

```
Out[187]: City                object
Colors Reported             object
Shape Reported              object
State                       object
Time                        datetime64[ns]
dtype: object
```

Documentation for `to_datetime` [.\(http://pandas.pydata.org/pandas-docs/stable/generated/pandas.to_datetime.html\)](http://pandas.pydata.org/pandas-docs/stable/generated/pandas.to_datetime.html)

```
In [188]: # convenient Series attributes are now available  
ufo.Time.dt.hour.head()
```

```
Out[188]: 0    22  
1    20  
2    14  
3    13  
4    19  
Name: Time, dtype: int64
```

```
In [189]: # ufo.Time.dt.weekday_name.head()
```

```
Out[189]: 0    Sunday  
1    Monday  
2    Sunday  
3    Monday  
4    Tuesday  
Name: Time, dtype: object
```

```
In [190]: # ufo.Time.dt.dayofyear.head()
```

```
Out[190]: 0    152  
1    181  
2     46  
3    152  
4    108  
Name: Time, dtype: int64
```

API reference for [datetimelike properties and methods](http://pandas.pydata.org/pandas-docs/stable/api.html#datetimelike-properties) (<http://pandas.pydata.org/pandas-docs/stable/api.html#datetimelike-properties>)

```
In [191]: # convert a single string to datetime format (outputs a timestamp object)  
ts = pd.to_datetime('1/1/1999')  
ts
```

```
Out[191]: Timestamp('1999-01-01 00:00:00')
```

```
In [192]: # compare a datetime Series with a timestamp  
ufo.loc[ufo.Time >= ts, :].head()
```

```
Out[192]:
```

	City	Colors Reported	Shape Reported	State	Time
12832	Loma Rica	NaN	LIGHT	CA	1999-01-01 02:30:00
12833	Bauxite	NaN	NaN	AR	1999-01-01 03:00:00
12834	Florence	NaN	CYLINDER	SC	1999-01-01 14:00:00
12835	Lake Henshaw	NaN	CIGAR	CA	1999-01-01 15:00:00
12836	Wilmington Island	NaN	LIGHT	GA	1999-01-01 17:15:00

```
In [193]: ▶ # perform mathematical operations with timestamps (outputs a timedelta object)
ufo.Time.max() - ufo.Time.min()
```

```
Out[193]: Timedelta('25781 days 01:59:00')
```

```
In [194]: ▶ # timedelta objects also have attributes you can access
(ufo.Time.max() - ufo.Time.min()).days
```

```
Out[194]: 25781
```


```
In [195]: ▶ # allow plots to appear in the notebook
%matplotlib inline
```

```
In [196]: ▶ # count the number of UFO reports per year
ufo['Year'] = ufo.Time.dt.year
ufo.Year.value_counts().sort_index().head()
```

```
Out[196]: 1930      2
          1931      2
          1933      1
          1934      1
          1935      1
          Name: Year, dtype: int64
```

```
In [197]: ▶ # plot the number of UFO reports per year (line plot is the default)
ufo.Year.value_counts().sort_index().plot()
```

```
Out[197]: <matplotlib.axes._subplots.AxesSubplot at 0x11ffb903b48>
```



How do I find and remove duplicate rows in pandas?

```
In [198]: # read a dataset of movie reviewers into a DataFrame
user_cols = ['user_id', 'age', 'gender', 'occupation', 'zip_code']
users = pd.read_table('http://bit.ly/movieusers', sep='|', header=None, names=user_cols)
users.head()
```

Out[198]:

	age	gender	occupation	zip_code
user_id				
1	24	M	technician	85711
2	53	F	other	94043
3	23	M	writer	32067
4	24	M	technician	43537
5	33	F	other	15213

```
In [199]: users.shape
```

Out[199]: (943, 4)

```
In [200]: # detect duplicate zip codes: True if an item is identical to a previous item
users.zip_code.duplicated().tail()
```

Out[200]:

user_id	
939	False
940	True
941	False
942	False
943	False

Name: zip_code, dtype: bool

```
In [201]: # count the duplicate items (True becomes 1, False becomes 0)
users.zip_code.duplicated().sum()
```

Out[201]: 148

```
In [202]: # detect duplicate DataFrame rows: True if an entire row is identical to a previous row
users.duplicated().tail()
```

Out[202]:

user_id	
939	False
940	False
941	False
942	False
943	False

dtype: bool


```
In [203]: # count the duplicate rows  
users.duplicated().sum()
```

Out[203]: 7

Logic for **duplicated** (<http://pandas.pydata.org/pandas-docs/stable/generated/pandas.DataFrame.duplicated.html>):

- **keep='first'** (default): Mark duplicates as True except for the first occurrence.
- **keep='last'** : Mark duplicates as True except for the last occurrence.
- **keep=False** : Mark all duplicates as True.

```
In [204]: # examine the duplicate rows (ignoring the first occurrence)  
users.loc[users.duplicated(keep='first'), :]
```

Out[204]:

	age	gender	occupation	zip_code
user_id				
496	21	F	student	55414
572	51	M	educator	20003
621	17	M	student	60402
684	28	M	student	55414
733	44	F	other	60630
805	27	F	other	20009
890	32	M	student	97301

```
In [205]: # examine the duplicate rows (ignoring the last occurrence)  
users.loc[users.duplicated(keep='last'), :]
```

Out[205]:

	age	gender	occupation	zip_code
user_id				
67	17	M	student	60402
85	51	M	educator	20003
198	21	F	student	55414
350	32	M	student	97301
428	28	M	student	55414
437	27	F	other	20009
460	44	F	other	60630

```
In [206]: # examine the duplicate rows (including all duplicates)  
users.loc[users.duplicated(keep=False), :]
```

Out[206]:

	age	gender	occupation	zip_code
user_id				
67	17	M	student	60402
85	51	M	educator	20003
198	21	F	student	55414
350	32	M	student	97301
428	28	M	student	55414
437	27	F	other	20009
460	44	F	other	60630
496	21	F	student	55414
572	51	M	educator	20003
621	17	M	student	60402
684	28	M	student	55414
733	44	F	other	60630
805	27	F	other	20009
890	32	M	student	97301

```
In [207]: # drop the duplicate rows (inplace=False by default)  
users.drop_duplicates(keep='first').shape
```

Out[207]: (936, 4)

```
In [208]: users.drop_duplicates(keep='last').shape
```

Out[208]: (936, 4)

```
In [209]: users.drop_duplicates(keep=False).shape
```

Out[209]: (929, 4)

Documentation for **drop_duplicates** (http://pandas.pydata.org/pandas-docs/stable/generated/pandas.DataFrame.drop_duplicates.html)

```
In [210]: # only consider a subset of columns when identifying duplicates  
users.duplicated(subset=['age', 'zip_code']).sum()
```

Out[210]: 16

```
In [211]: users.drop_duplicates(subset=['age', 'zip_code']).shape
```

Out[211]: (927, 4)

How do I avoid a SettingWithCopyWarning in pandas?

```
In [212]: # read a dataset of top-rated IMDb movies into a DataFrame  
movies = pd.read_csv('http://bit.ly/imdbratings')  
movies.head()
```

Out[212]:

	star_rating	title	content_rating	genre	duration	actors_list
0	9.3	The Shawshank Redemption	R	Crime	142	[u'Tim Robbins', u'Morgan Freeman', u'Bob Gunt...
1	9.2	The Godfather	R	Crime	175	[u'Marlon Brando', u'Al Pacino', u'James Caan']
2	9.1	The Godfather: Part II	R	Crime	200	[u'Al Pacino', u'Robert De Niro', u'Robert Duv...
3	9.0	The Dark Knight	PG-13	Action	152	[u'Christian Bale', u'Heath Ledger', u'Aaron E...
4	8.9	Pulp Fiction	R	Crime	154	[u'John Travolta', u'Uma Thurman', u'Samuel L....

```
In [213]: # count the missing values in the 'content_rating' Series  
movies.content_rating.isnull().sum()
```

Out[213]: 3

```
In [214]: # examine the DataFrame rows that contain those missing values  
movies[movies.content_rating.isnull()]
```

Out[214]:

	star_rating	title	content_rating	genre	duration	actors_list
187	8.2	Butch Cassidy and the Sundance Kid	NaN	Biography	110	[u'Paul Newman', u'Robert Redford', u'Katharin...
649	7.7	Where Eagles Dare	NaN	Action	158	[u'Richard Burton', u'Clint Eastwood', u'Mary ...
936	7.4	True Grit	NaN	Adventure	128	[u'John Wayne', u'Kim Darby', u'Glen Campbell']

```
In [215]: # examine the unique values in the 'content_rating' Series
movies.content_rating.value_counts()
```

```
Out[215]: R          460
PG-13       189
PG          123
NOT RATED   65
APPROVED    47
UNRATED     38
G           32
NC-17       7
PASSED      7
X           4
GP          3
TV-MA       1
Name: content_rating, dtype: int64
```

Goal: Mark the 'NOT RATED' values as missing values, represented by 'NaN'.

```
In [216]: # first, locate the relevant rows
movies[movies.content_rating=='NOT RATED'].head()
```

```
Out[216]:
```

	star_rating	title	content_rating	genre	duration	actors_list
5	8.9	12 Angry Men	NOT RATED	Drama	96	[u'Henry Fonda', u'Lee J. Cobb', u'Martin Bals...
6	8.9	The Good, the Bad and the Ugly	NOT RATED	Western	161	[u'Clint Eastwood', u'Eli Wallach', u'Lee Van ...
41	8.5	Sunset Blvd.	NOT RATED	Drama	110	[u'William Holden', u'Gloria Swanson', u'Erich...
63	8.4	M	NOT RATED	Crime	99	[u'Peter Lorre', u'Ellen Widmann', u'Inge Land...
66	8.4	Munna Bhai M.B.B.S.	NOT RATED	Comedy	156	[u'Sunil Dutt', u'Sanjay Dutt', u'Arshad Warsi']

```
In [217]: # then, select the 'content_rating' Series from those rows
movies[movies.content_rating=='NOT RATED'].content_rating.head()
```

```
Out[217]: 5    NOT RATED
6    NOT RATED
41   NOT RATED
63   NOT RATED
66   NOT RATED
Name: content_rating, dtype: object
```

```
In [279]: # finally, replace the 'NOT RATED' values with 'NaN' (imported from NumPy)
import numpy as np
movies[movies.content_rating=='NOT RATED'].content_rating = np.nan
```

Problem: That statement involves two operations, a `__getitem__` and a `__setitem__`. pandas can't guarantee whether the `__getitem__` operation returns a view or a copy of the data.

- If `__getitem__` returns a view of the data, `__setitem__` will affect the 'movies' DataFrame.
- But if `__getitem__` returns a copy of the data, `__setitem__` will not affect the 'movies' DataFrame.

```
In [219]: # the 'content_rating' Series has not changed
movies.content_rating.isnull().sum()
```

Out[219]: 3

Solution: Use the `loc` method, which replaces the 'NOT RATED' values in a single `__setitem__` operation.

```
In [220]: # replace the 'NOT RATED' values with 'NaN' (does not cause a SettingWithCopy)
movies.loc[movies.content_rating=='NOT RATED', 'content_rating'] = np.nan
```

```
In [221]: # this time, the 'content_rating' Series has changed
movies.content_rating.isnull().sum()
```

Out[221]: 68

Summary: Use the `loc` method any time you are selecting rows and columns in the same statement.

More information: [Modern Pandas \(Part 1\)](http://tomaugspurger.github.io/modern-1.html) (<http://tomaugspurger.github.io/modern-1.html>).

```
In [222]: # create a DataFrame only containing movies with a high 'star_rating'
top_movies = movies.loc[movies.star_rating >= 9, :]
top_movies
```

Out[222]:

	star_rating	title	content_rating	genre	duration	actors_list
0	9.3	The Shawshank Redemption	R	Crime	142	[u'Tim Robbins', u'Morgan Freeman', u'Bob Gunt...]
1	9.2	The Godfather	R	Crime	175	[u'Marlon Brando', u'Al Pacino', u'James Caan']
2	9.1	The Godfather: Part II	R	Crime	200	[u'Al Pacino', u'Robert De Niro', u'Robert Duv...]
3	9.0	The Dark Knight	PG-13	Action	152	[u'Christian Bale', u'Heath Ledger', u'Aaron E...]

Goal: Fix the 'duration' for 'The Shawshank Redemption'.

```
In [276]: ▶ # overwrite the relevant cell with the correct duration
top_movies.loc[0, 'duration'] = 150
```

Problem: pandas isn't sure whether 'top_movies' is a view or a copy of 'movies'.

```
In [275]: ▶ # 'top_movies' DataFrame has been updated
top_movies
```

Out[275]:

	star_rating	title	content_rating	genre	duration	actors_list	4
0	9.3	The Shawshank Redemption	R	Crime	150	[u'Tim Robbins', u'Morgan Freeman', u'Bob Gunt...	150.0
1	9.2	The Godfather	R	Crime	175	[u'Marlon Brando', u'Al Pacino', u'James Caan']	NaN
2	9.1	The Godfather: Part II	R	Crime	200	[u'Al Pacino', u'Robert De Niro', u'Robert Duv...	NaN
3	9.0	The Dark Knight	PG-13	Action	152	[u'Christian Bale', u'Heath Ledger', u'Aaron E...	NaN

```
In [225]: ▶ # 'movies' DataFrame has not been updated
movies.head(1)
```

Out[225]:

	star_rating	title	content_rating	genre	duration	actors_list
0	9.3	The Shawshank Redemption	R	Crime	142	[u'Tim Robbins', u'Morgan Freeman', u'Bob Gunt...

Solution: Any time you are attempting to create a DataFrame copy, use the [copy](http://pandas.pydata.org/pandas-docs/stable/generated/pandas.DataFrame.copy.html) (<http://pandas.pydata.org/pandas-docs/stable/generated/pandas.DataFrame.copy.html>) method.

```
In [226]: ▶ # explicitly create a copy of 'movies'
top_movies = movies.loc[movies.star_rating >= 9, :].copy()
```

```
In [227]: ▶ # pandas now knows that you are updating a copy instead of a view (does not c
top_movies.loc[0, 'duration'] = 150
```

```
In [228]: ▶ # 'top_movies' DataFrame has been updated
top_movies
```

Out[228]:

	star_rating	title	content_rating	genre	duration	actors_list
0	9.3	The Shawshank Redemption	R	Crime	150	[u'Tim Robbins', u'Morgan Freeman', u'Bob Gunt...
1	9.2	The Godfather	R	Crime	175	[u'Marlon Brando', u'Al Pacino', u'James Caan']
2	9.1	The Godfather: Part II	R	Crime	200	[u'Al Pacino', u'Robert De Niro', u'Robert Duv...
3	9.0	The Dark Knight	PG-13	Action	152	[u'Christian Bale', u'Heath Ledger', u'Aaron E...

Documentation on indexing and selection: [Returning a view versus a copy](http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy).
(<http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>)

Stack Overflow: [What is the point of views in pandas if it is undefined whether an indexing operation returns a view or a copy?](http://stackoverflow.com/questions/34884536/what-is-the-point-of-views-in-pandas-if-it-is-undefined-whether-an-indexing-operation-returns-a-view-or-a-copy?) (<http://stackoverflow.com/questions/34884536/what-is-the-point-of-views-in-pandas-if-it-is-undefined-whether-an-indexing-oper>)

How do I change display options in pandas?

```
In [229]: ▶ # read a dataset of alcohol consumption into a DataFrame
drinks = pd.read_csv('http://bit.ly/drinksbycountry')
```

```
In [230]: ▶ # only 60 rows will be displayed when printing
drinks
```

Out[230]:

	country	beer_servings	spirit_servings	wine_servings	total_litres_of_pure_alcohol
0	Afghanistan	0	0	0	0.0
1	Albania	89	132	54	4.9
2	Algeria	25	0	14	0.7
3	Andorra	245	138	312	12.4
4	Angola	217	57	45	5.9
...
188	Venezuela	333	100	3	7.7
189	Vietnam	111	2	1	2.0
190	Yemen	6	0	0	0.1
191	Zambia	32	19	4	2.5
192	Zimbabwe	64	18	4	4.7

193 rows × 6 columns

```
In [231]: ▶ # check the current setting for the 'max_rows' option
pd.get_option('display.max_rows')
```

Out[231]: 60

Documentation for **get_option** [_\(http://pandas.pydata.org/pandas-docs/stable/generated/pandas.get_option.html\)](http://pandas.pydata.org/pandas-docs/stable/generated/pandas.get_option.html)


```
In [232]: # overwrite the current setting so that all rows will be displayed  
pd.set_option('display.max_rows', None)  
drinks
```

Out[232]:

	country	beer_servings	spirit_servings	wine_servings	total_litres_of_pure_alcohol
0	Afghanistan	0	0	0	0.0
1	Albania	89	132	54	4.9
2	Algeria	25	0	14	0.7
3	Andorra	245	138	312	12.4
4	Angola	217	57	45	5.9
5	Antigua & Barbuda	102	128	45	4.9
6	Argentina	193	25	221	8.3
7	Armenia	21	179	11	3.8
8	Australia	261	72	212	10.4

```
In [233]: # reset the 'max_rows' option to its default  
pd.reset_option('display.max_rows')
```

Documentation for **set_option** (http://pandas.pydata.org/pandas-docs/stable/generated/pandas.set_option.html) and **reset_option** (http://pandas.pydata.org/pandas-docs/stable/generated/pandas.reset_option.html)

```
In [234]: # the 'max_columns' option is similar to 'max_rows'  
pd.get_option('display.max_columns')
```

Out[234]: 20


```
In [235]: # read the training dataset from Kaggle's Titanic competition into a DataFrame
train = pd.read_csv('http://bit.ly/kaggletrain')
train.head()
```

Out[235]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	71.2833
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500

```
In [236]: # an ellipsis is displayed in the 'Name' cell of row 1 because of the 'max_colwidth'
pd.get_option('display.max_colwidth')
```

Out[236]: 50

```
In [237]:  # overwrite the current setting so that more characters will be displayed
pd.set_option('display.max_colwidth', 1000)
train.head()
```

Out[237]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Thayer)	female	38.0	1	0	PC 17599	71.2833
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500

```
In [238]: # overwrite the 'precision' setting to display 2 digits after the decimal point
pd.set_option('display.precision', 2)
train.head()
```

Out[238]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	C
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.25	
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Thayer)	female	38.0	1	0	PC 17599	71.28	
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.92	
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.10	C
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.05	

```
In [239]: # add two meaningless columns to the drinks DataFrame
drinks['x'] = drinks.wine_servings * 1000
drinks['y'] = drinks.total_litres_of_pure_alcohol * 1000
drinks.head()
```

Out[239]:

	country	beer_servings	spirit_servings	wine_servings	total_litres_of_pure_alcohol	continent
0	Afghanistan	0	0	0	0.0	
1	Albania	89	132	54	4.9	Eu
2	Algeria	25	0	14	0.7	A
3	Andorra	245	138	312	12.4	Eu
4	Angola	217	57	45	5.9	A

```
In [240]: # use a Python format string to specify a comma as the thousands separator
pd.set_option('display.float_format', '{:,}'.format)
drinks.head()
```

Out[240]:

	country	beer_servings	spirit_servings	wine_servings	total_litres_of_pure_alcohol	continent
0	Afghanistan	0	0	0	0.0	
1	Albania	89	132	54	4.9	Eu
2	Algeria	25	0	14	0.7	A
3	Andorra	245	138	312	12.4	Eu
4	Angola	217	57	45	5.9	A

```
In [241]: # 'y' was affected (but not 'x') because the 'float_format' option only affected
drinks.dtypes
```

```
Out[241]: country          object
beer_servings          int64
spirit_servings        int64
wine_servings          int64
total_litres_of_pure_alcohol  float64
continent              object
x                      int64
y                      float64
dtype: object
```

```
In [242]: # view the option descriptions (including the default and current values)
pd.describe_option()
```

```
Whether to use the Unicode East Asian Width to calculate the display
text
width.
Enabling this may affect to the performance (default: False)
[default: False] [currently: False]display.unicode.east_asian_width :
boolean
Whether to use the Unicode East Asian Width to calculate the display
text
width.
Enabling this may affect to the performance (default: False)
[default: False] [currently: False]display.width : int
Width of the display in characters. In case python/IPython is running
in
a terminal this can be set to None and pandas will correctly auto-det
ect
the width.
Note that the IPython notebook, IPython qtconsole, or IDLE do not run
in a
terminal and hence it is not possible to correctly detect the width.
[default: 80] [currently: 80]io.excel.ods.reader : string
```

```
In [243]: ▶ # search for specific options by name
pd.describe_option('rows')
```

display.max_info_rows : int or None
df.info() will usually show null-counts for each column.
For large frames this can be quite slow. max_info_rows and max_info_col
s

limit this null check only to frames with smaller dimensions than specified.
[default: 1690785] [currently: 1690785]display.max_rows : int
If max_rows is exceeded, switch to truncate view. Depending on `large_repr`, objects are either centrally truncated or printed as a summary view. 'None' value means unlimited.

In case python/IPython is running in a terminal and `large_repr` equals 'truncate' this can be set to 0 and pandas will auto-detect the height of the terminal and print a truncated object which fits the screen height. The IPython notebook, IPython qtconsole, or IDLE do not run in a terminal and hence it is not possible to do correct auto-detection.

[default: 60] [currently: 60]display.min_rows : int
The numbers of rows to show in a truncated view (when `max_rows` is exceeded). Ignored when `max_rows` is set to None or 0. When set to None, follows the value of `max_rows`.
[default: 10] [currently: 10]

Documentation for **describe_option** (http://pandas.pydata.org/pandas-docs/stable/generated/pandas.describe_option.html)

```
In [277]: ▶ # reset all of the options to their default values
pd.reset_option('all')
```

: boolean
use_inf_as_null had been deprecated and will be removed in a future version. Use `use_inf_as_na` instead.

How do I create a pandas DataFrame from another object?

```
In [245]: ▶ # create a DataFrame from a dictionary (keys become column names, values become data)
pd.DataFrame({'id':[100, 101, 102], 'color':['red', 'blue', 'red']})
```

Out[245]:

	id	color
0	100	red
1	101	blue
2	102	red

```
In [246]: ▶ # optionally specify the order of columns and define the index
df = pd.DataFrame({'id':[100, 101, 102], 'color':['red', 'blue', 'red']}, col
df
```

Out[246]:

	id	color
a	100	red
b	101	blue
c	102	red

Documentation for **DataFrame** (<http://pandas.pydata.org/pandas-docs/stable/generated/pandas.DataFrame.html>)

```
In [247]: ▶ # create a DataFrame from a list of lists (each inner list becomes a row)
pd.DataFrame([[100, 'red'], [101, 'blue'], [102, 'red']], columns=['id', 'col
```

Out[247]:

	id	color
0	100	red
1	101	blue
2	102	red

```
In [248]: ▶ # create a NumPy array (with shape 4 by 2) and fill it with random numbers be
import numpy as np
arr = np.random.rand(4, 2)
arr
```

Out[248]: array([[0.76905988, 0.03304819],
[0.0732601 , 0.32422705],
[0.10660698, 0.16173659],
[0.65393043, 0.02880436]])

```
In [249]: ▶ # create a DataFrame from the NumPy array
pd.DataFrame(arr, columns=['one', 'two'])
```

Out[249]:

	one	two
0	0.769060	0.033048
1	0.073260	0.324227
2	0.106607	0.161737
3	0.653930	0.028804

```
In [250]: ▶ # create a DataFrame of student IDs (100 through 109) and test scores (random
pd.DataFrame({'student':np.arange(100, 110, 1), 'test':np.random.randint(60,
```

Out[250]:

	student	test
0	100	64
1	101	83
2	102	92
3	103	100
4	104	60
5	105	67
6	106	82
7	107	79
8	108	80
9	109	66

Documentation for [np.arange](http://docs.scipy.org/doc/numpy/reference/generated/numpy.arange.html)

(<http://docs.scipy.org/doc/numpy/reference/generated/numpy.arange.html>) and [np.random](http://docs.scipy.org/doc/numpy/reference/routines.random.html)

(<http://docs.scipy.org/doc/numpy/reference/routines.random.html>)

```
In [251]: ▶ # 'set_index' can be chained with the DataFrame constructor to select an index
pd.DataFrame({'student':np.arange(100, 110, 1), 'test':np.random.randint(60,
```

Out[251]:

	test
student	
100	83
101	82
102	94
103	96
104	69
105	72
106	98
107	73
108	95
109	90

Documentation for [set_index](http://pandas.pydata.org/pandas-docs/stable/generated/pandas.DataFrame.set_index.html) (http://pandas.pydata.org/pandas-docs/stable/generated/pandas.DataFrame.set_index.html)


```
In [252]: # create a new Series using the Series constructor  
s = pd.Series(['round', 'square'], index=['c', 'b'], name='shape')  
s
```

```
Out[252]: c    round  
         b    square  
         Name: shape, dtype: object
```

Documentation for **Series** (<http://pandas.pydata.org/pandas-docs/stable/generated/pandas.Series.html>)

```
In [271]: # concatenate the DataFrame and the Series (use axis=1 to concatenate columns)  
pd.concat([df, s], axis=1, sort=True)
```

```
Out[271]:
```

	id	color	shape
a	100	red	NaN
b	101	blue	square
c	102	red	round

Notes:

- The Series name became the column name in the DataFrame.
- The Series data was aligned to the DataFrame by its index.
- The 'shape' for row 'a' was marked as a missing value (NaN) because that index was not present in the Series.

Documentation for **concat** (<http://pandas.pydata.org/pandas-docs/stable/generated/pandas.concat.html>)

How do I apply a function to a pandas Series or DataFrame?

```
In [254]: # read the training dataset from Kaggle's Titanic competition into a DataFrame
train = pd.read_csv('http://bit.ly/kaggletrain')
train.head()
```

Out[254]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	71.2833
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500

Goal: Map the existing values of a Series to a different set of values

Method: `map` (<http://pandas.pydata.org/pandas-docs/stable/generated/pandas.Series.map.html>)
(Series method)

```
In [255]: # map 'female' to 0 and 'male' to 1
train['Sex_num'] = train.Sex.map({'female':0, 'male':1})
train.loc[0:4, ['Sex', 'Sex_num']]
```

Out[255]:

	Sex	Sex_num
0	male	1
1	female	0
2	female	0
3	female	0
4	male	1

Goal: Apply a function to each element in a Series

Method: `apply` (<http://pandas.pydata.org/pandas-docs/stable/generated/pandas.Series.apply.html>) (Series method)

Note: `map` can be substituted for `apply` in many cases, but `apply` is more flexible and thus is recommended

```
In [256]: ▶ # calculate the length of each string in the 'Name' Series
train['Name_length'] = train.Name.apply(len)
train.loc[0:4, ['Name', 'Name_length']]
```

Out[256]:

	Name	Name_length
0	Braund, Mr. Owen Harris	23
1	Cumings, Mrs. John Bradley (Florence Briggs Th...	51
2	Heikkinen, Miss. Laina	22
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	44
4	Allen, Mr. William Henry	24

```
In [257]: ▶ # round up each element in the 'Fare' Series to the next integer
import numpy as np
train['Fare_ceil'] = train.Fare.apply(np.ceil)
train.loc[0:4, ['Fare', 'Fare_ceil']]
```

Out[257]:

	Fare	Fare_ceil
0	7.2500	8.0
1	71.2833	72.0
2	7.9250	8.0
3	53.1000	54.0
4	8.0500	9.0

```
In [258]: ▶ # we want to extract the last name of each person
train.Name.head()
```

Out[258]:

0	Braund, Mr. Owen Harris
1	Cumings, Mrs. John Bradley (Florence Briggs Th...
2	Heikkinen, Miss. Laina
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)
4	Allen, Mr. William Henry

Name: Name, dtype: object

```
In [259]: ▶ # use a string method to split the 'Name' Series at commas (returns a Series
train.Name.str.split(',').head()
```

```
Out[259]: 0      [Braund, Mr. Owen Harris]
1      [Cumings, Mrs. John Bradley (Florence Briggs ...
2      [Heikkinen, Miss. Laina]
3      [Futrelle, Mrs. Jacques Heath (Lily May Peel)]
4      [Allen, Mr. William Henry]
Name: Name, dtype: object
```

```
In [260]: ▶ # define a function that returns an element from a list based on position
def get_element(my_list, position):
    return my_list[position]
```

```
In [261]: ▶ # apply the 'get_element' function and pass 'position' as a keyword argument
train.Name.str.split(',').apply(get_element, position=0).head()
```

```
Out[261]: 0      Braund
1      Cumings
2      Heikkinen
3      Futrelle
4      Allen
Name: Name, dtype: object
```

```
In [262]: ▶ # alternatively, use a lambda function
train.Name.str.split(',').apply(lambda x: x[0]).head()
```

```
Out[262]: 0      Braund
1      Cumings
2      Heikkinen
3      Futrelle
4      Allen
Name: Name, dtype: object
```

Goal: Apply a function along either axis of a DataFrame

Method: `apply` (<http://pandas.pydata.org/pandas-docs/stable/generated/pandas.DataFrame.apply.html>) (DataFrame method)

```
In [263]: # read a dataset of alcohol consumption into a DataFrame
drinks = pd.read_csv('http://bit.ly/drinksbycountry')
drinks.head()
```

Out[263]:

	country	beer_servings	spirit_servings	wine_servings	total_litres_of_pure_alcohol	continent
0	Afghanistan	0	0	0	0.0	
1	Albania	89	132	54	4.9	Europe
2	Algeria	25	0	14	0.7	Asia
3	Andorra	245	138	312	12.4	Europe
4	Angola	217	57	45	5.9	Asia

```
In [264]: # select a subset of the DataFrame to work with
drinks.loc[:, 'beer_servings':'wine_servings'].head()
```

Out[264]:

	beer_servings	spirit_servings	wine_servings
0	0	0	0
1	89	132	54
2	25	0	14
3	245	138	312
4	217	57	45

```
In [265]: # apply the 'max' function along axis 0 to calculate the maximum value in each column
drinks.loc[:, 'beer_servings':'wine_servings'].apply(max, axis=0)
```

```
Out[265]: beer_servings    376
           spirit_servings    438
           wine_servings    370
           dtype: int64
```

```
In [266]: # apply the 'max' function along axis 1 to calculate the maximum value in each row
drinks.loc[:, 'beer_servings':'wine_servings'].apply(max, axis=1).head()
```

```
Out[266]: 0      0
           1    132
           2     25
           3    312
           4    217
           dtype: int64
```

```
In [278]: ▶ # use 'np.argmax' to calculate which column has the maximum value for each row
drinks.loc[:, 'beer_servings':'wine_servings'].apply(np.argmax, axis=1).head()
```

```
Out[278]: 0    beer_servings
1    spirit_servings
2    beer_servings
3    wine_servings
4    beer_servings
dtype: object
```

Goal: Apply a function to every element in a DataFrame

Method: `applymap` (<http://pandas.pydata.org/pandas-docs/stable/generated/pandas.DataFrame.applymap.html>) (DataFrame method)

```
In [268]: ▶ # convert every DataFrame element into a float
drinks.loc[:, 'beer_servings':'wine_servings'].applymap(float).head()
```

```
Out[268]:
```

	beer_servings	spirit_servings	wine_servings
0	0.0	0.0	0.0
1	89.0	132.0	54.0
2	25.0	0.0	14.0
3	245.0	138.0	312.0
4	217.0	57.0	45.0

```
In [269]: ▶ # overwrite the existing DataFrame columns
drinks.loc[:, 'beer_servings':'wine_servings'] = drinks.loc[:, 'beer_servings':
drinks.head()
```

```
Out[269]:
```

	country	beer_servings	spirit_servings	wine_servings	total_litres_of_pure_alcohol	continent
0	Afghanistan	0.0	0.0	0.0	0.0	
1	Albania	89.0	132.0	54.0	4.9	Eu
2	Algeria	25.0	0.0	14.0	0.7	A
3	Andorra	245.0	138.0	312.0	12.4	Eu
4	Angola	217.0	57.0	45.0	5.9	A